

```

# import libraries
import warnings
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost as xgb
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_curve, auc, precision_score, recall_score,
f1_score, roc_auc_score

#load the data
data = pd.read_csv('/content/framingham_heart_study.csv')
data.head()

{"summary": "{\n    \"name\": \"data\", \n    \"rows\": 4240, \n    \"fields\": [\n        {\n            \"column\": \"male\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0, \n                \"min\": 0, \n                \"max\": 1, \n                \"num_unique_values\": 2, \n                \"samples\": [\n                    0, \n                    1\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }, \n        {\n            \"column\": \"age\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 8, \n                \"min\": 32, \n                \"max\": 70, \n                \"num_unique_values\": 39, \n                \"samples\": [\n                    34, \n                    70\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }, \n        {\n            \"column\": \"education\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 1.0197911793650334, \n                \"min\": 1.0, \n                \"max\": 4.0, \n                \"num_unique_values\": 4, \n                \"samples\": [\n                    2.0, \n                    3.0\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }, \n        {\n            \"column\": \"currentSmoker\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0, \n                \"min\": 0, \n                \"max\": 1, \n                \"num_unique_values\": 2, \n                \"samples\": [\n                    1, \n                    0\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }, \n        {\n            \"column\": \"cigsPerDay\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 11.922461800608747, \n                \"min\": 0.0, \n                \"max\": 70.0, \n                \"num_unique_values\": 33, \n                \"samples\": [\n                    19.0, \n                    4.0\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }, \n        {\n            \"column\": \"BPMeds\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 0.16954428739625657, \n                \"min\": 0.0, \n                \"max\": 1.0, \n                \"samples\": [\n                    0.0, \n                    1.0
                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n            }\\n        }
    ]
}

```

```

  "num_unique_values": 2,\n      "samples": [\n          1.0,\n0.0\n      ],\n      "semantic_type": "\\",\\n\n  "description": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"prevalentStroke\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 0,\n              \"min\": 0,\n              \"max\": 1,\n      \"num_unique_values\": 2,\n      \"samples\": [\n          1,\n              0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"prevalentHyp\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 0,\n              \"min\": 0,\n              \"max\": 1,\n      \"num_unique_values\": 2,\n      \"samples\": [\n          1,\n              0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"diabetes\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 0,\n              \"min\": 0,\n              \"max\": 1,\n      \"num_unique_values\": 2,\n      \"samples\": [\n          1,\n              0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"totChol\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 44.59128386860702,\n              \"min\":\n107.0,\n              \"max\": 696.0,\n      \"num_unique_values\": 248,\n      \"samples\": [\n          311.0,\n              205.0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"sysBP\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 22.0332996088492,\n              \"min\": 83.5,\n              \"max\": 295.0,\n      \"num_unique_values\": 234,\n      \"samples\": [\n          109.0,\n              184.5\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"diaBP\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 11.910394483305936,\n              \"min\": 48.0,\n              \"max\": 142.5,\n      \"num_unique_values\": 146,\n      \"samples\": [\n          106.0,\n              108.5\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"BMI\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 4.079840168944382,\n              \"min\":\n15.54,\n              \"max\": 56.8,\n      \"num_unique_values\": 1364,\n      \"samples\": [\n          24.56,\n              19.87\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"heartRate\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 12.025347984469342,\n              \"min\": 44.0,\n              \"max\": 143.0,\n      \"num_unique_values\": 73,\n      \"samples\": [\n          85.0,\n              47.0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  },\\n  {\n      \"column\": \n  \"glucose\",\\n      \"properties\": {\n          \"number\":,\n              \"std\": 23.95433481134474,\n              \"min\": 40.0,\n              \"max\": 394.0,\n      \"num_unique_values\": 143,\n      \"samples\": [\n          74.0\\n              394.0\\n      ],\n      \"semantic_type\": \n  \"\",\\n\n  \"description\": \"\"\\n      }\\n  }\n}
```

```

  "semantic_type": "\",
    "description": """
      },
      {
        "column": "TenYearCHD",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 0,
          "max": 1,
          "num_unique_values": 2,
          "samples": [
            0
          ],
          "semantic_type": "\",
          "description": """
        }
      }
    }
  },
  "type": "dataframe",
  "variable_name": "data"
}

# Preview all columns
data.columns

Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay',
       'BPMed', 'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol',
       'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'],
      dtype='object')

# Preview counts of NULL values
data.isnull().sum()

male          0
age           0
education     105
currentSmoker 0
cigsPerDay    29
BPMed         53
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        50
sysBP          0
diaBP          0
BMI            19
heartRate      1
glucose        388
TenYearCHD     0
dtype: int64

# Preview counts of duplicated records
data.duplicated().sum()

np.int64(0)

# Preview dataset shape; rows: columns
data.shape

(4240, 16)

```

```
# Drop records with NULL values
data.dropna(inplace=True)

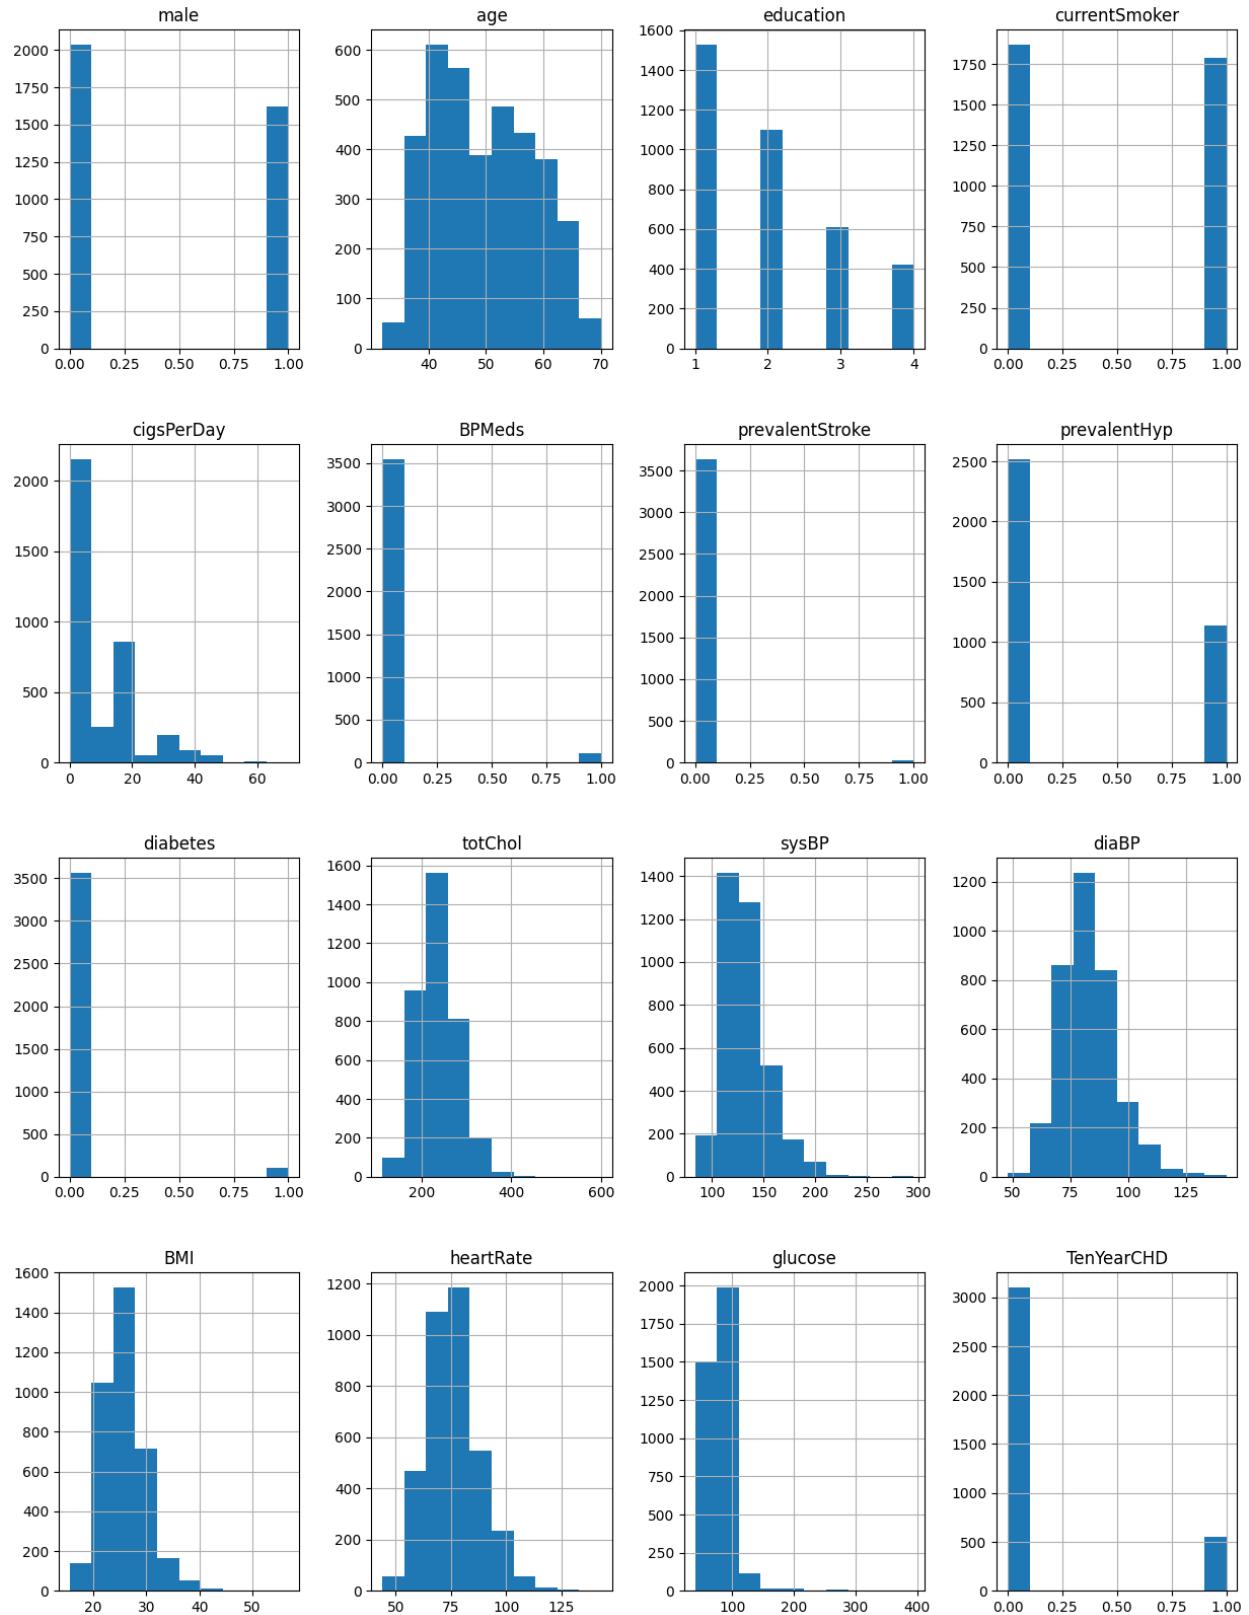
# Verify that there are no NULL values left
data.isnull().sum()

male          0
age           0
education     0
currentSmoker 0
cigsPerDay    0
BPMed          0
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        0
sysBP          0
diaBP          0
BMI            0
heartRate      0
glucose         0
TenYearCHD     0
dtype: int64
```

Data Distribution

```
# plot histogram to see the distribution of the data
fig = plt.figure(figsize = (15,20))
ax = fig.gca()
data.hist(ax = ax)
plt.show()

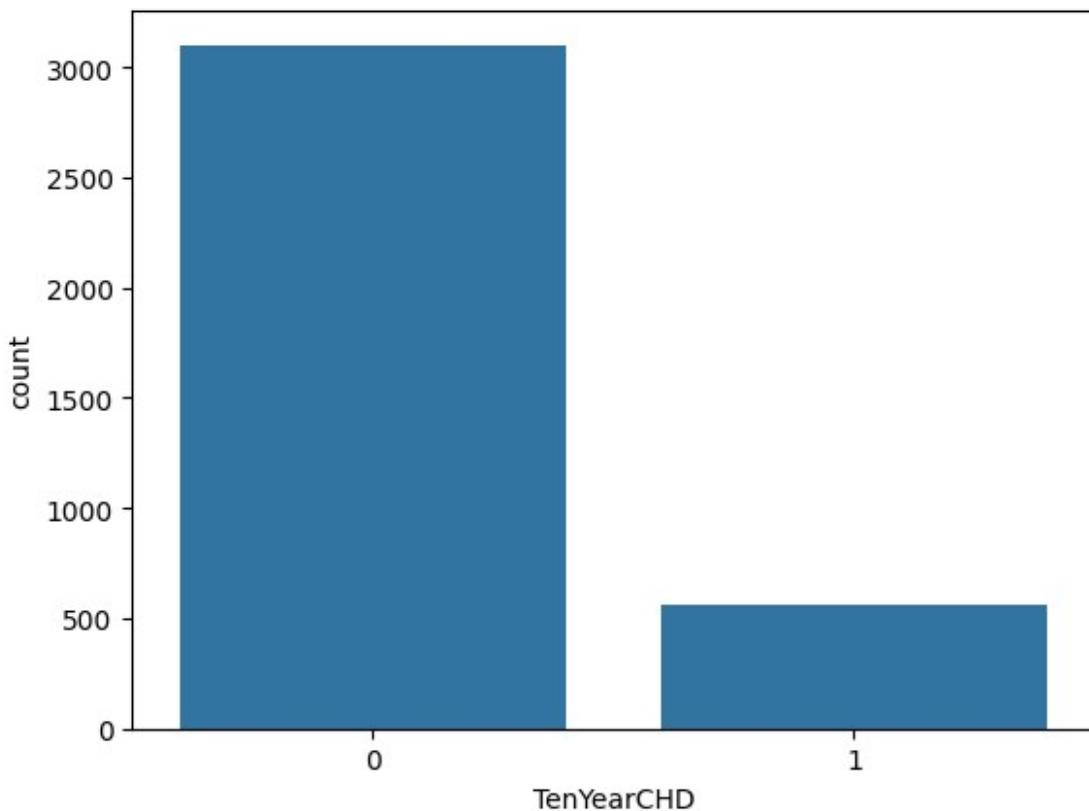
/tmp/ipython-input-887923207.py:4: UserWarning: To output multiple
subplots, the figure containing the passed axes is being cleared.
  data.hist(ax = ax)
```



```

sns.countplot(x='TenYearCHD', data=data)
plt.show()
cases = data.TenYearCHD.value_counts()
print(f"There are {cases[0]} patients without heart disease and
{cases[1]} patients with the disease")

```



There are 3101 patients without heart disease and 557 patients with the disease

The data is not properly balanced as the number of people without the disease greatly exceeds the number of people with the disease. The ratio is about 1:5.57

```

from operator import add

def stacked_barchart(data, title = None, ylabel = None, xlabel = None):
    default_colors = ['#008080', '#5f3c41', '#219AD8']
    # From raw value to percentage
    totals = data.sum(axis=1)
    bars = ((data.T / totals) * 100).T
    r = list(range(data.index.size))

    # Plot

```

```

barWidth = 0.95
names = data.index.tolist()
bottom = [0] * bars.shape[0]

# Create bars
color_index = 0
plots = []
for bar in bars.columns:
    plots.append(plt.bar(r, bars[bar], bottom=bottom,
color=default_colors[color_index], edgecolor='white', width=barWidth))
    bottom = list(map(add, bottom, bars[bar]))
    color_index = 0 if color_index >= len(default_colors) else
color_index + 1

# Custom x axis
plt.title(title)
plt.xticks(r, names)
plt.xlabel(data.index.name if xlabel is None else xlabel)
plt.ylabel(data.columns.name if ylabel is None else ylabel)
ax = plt.gca()

y_labels = ax.get_yticks()
ax.set_yticklabels([str(y) + '%' for y in y_labels])

flat_list = [item for sublist in data.T.values for item in
sublist]
for i, d in zip(ax.patches, flat_list):
    data_label = str(d) + " (" + str(round(i.get_height(), 2)) +
"%)"
    ax.text(i.get_x() + 0.45, i.get_y() + 5, data_label,
horizontalalignment='center', verticalalignment='center', fontdict =
dict(color = 'white', size = 20))

for item in ([ax.title]):
    item.set_fontsize(27)

for item in ([ax.xaxis.label, ax.yaxis.label] +
ax.get_xticklabels() + ax.get_yticklabels()):
    item.set_fontsize(24)

legend = ax.legend(plots, bars.columns.tolist(), fancybox=True)
plt.setp(legend.get_texts(), fontsize='20')

fig = plt.gcf()
fig.set_size_inches(25, 35)
grid_rows = 3
grid_cols = 2

#draw sex vs disease outcome
plt.subplot(grid_rows, grid_cols, 1)

```

```

temp =
data[['male','TenYearCHD']].groupby(['male','TenYearCHD']).size().unstack('TenYearCHD')
temp.rename(index={0:'Female', 1:'Male'}, columns={0:'No Disease',
1:'Has Disease'}, inplace = True)
stacked_barchart(temp, title = 'CHD vs Sex', ylabel = 'Population')

#draw smoking status vs disease outcome
plt.subplot(grid_rows, grid_cols, 2)
temp =
data[['currentSmoker','TenYearCHD']].groupby(['currentSmoker','TenYear
CHD']).size().unstack('TenYearCHD')
temp.rename(index={0:'Not a Smoker', 1:'Smoker'}, columns={0:'No
Disease', 1:'Has Disease'}, inplace = True)
stacked_barchart(temp, title = 'CHD vs Smoking', ylabel =
'Population')

#draw diabetes vs disease outcome
plt.subplot(grid_rows, grid_cols, 3)
temp =
data[['diabetes','TenYearCHD']].groupby(['diabetes','TenYearCHD']).size()
.unstack('TenYearCHD')
temp.rename(index={0:'Not Diabetic', 1:'Diabetic'}, columns={0:'No
Disease', 1:'Has Disease'}, inplace = True)
stacked_barchart(temp, title = 'CHD vs Diabetes', ylabel =
'Population')

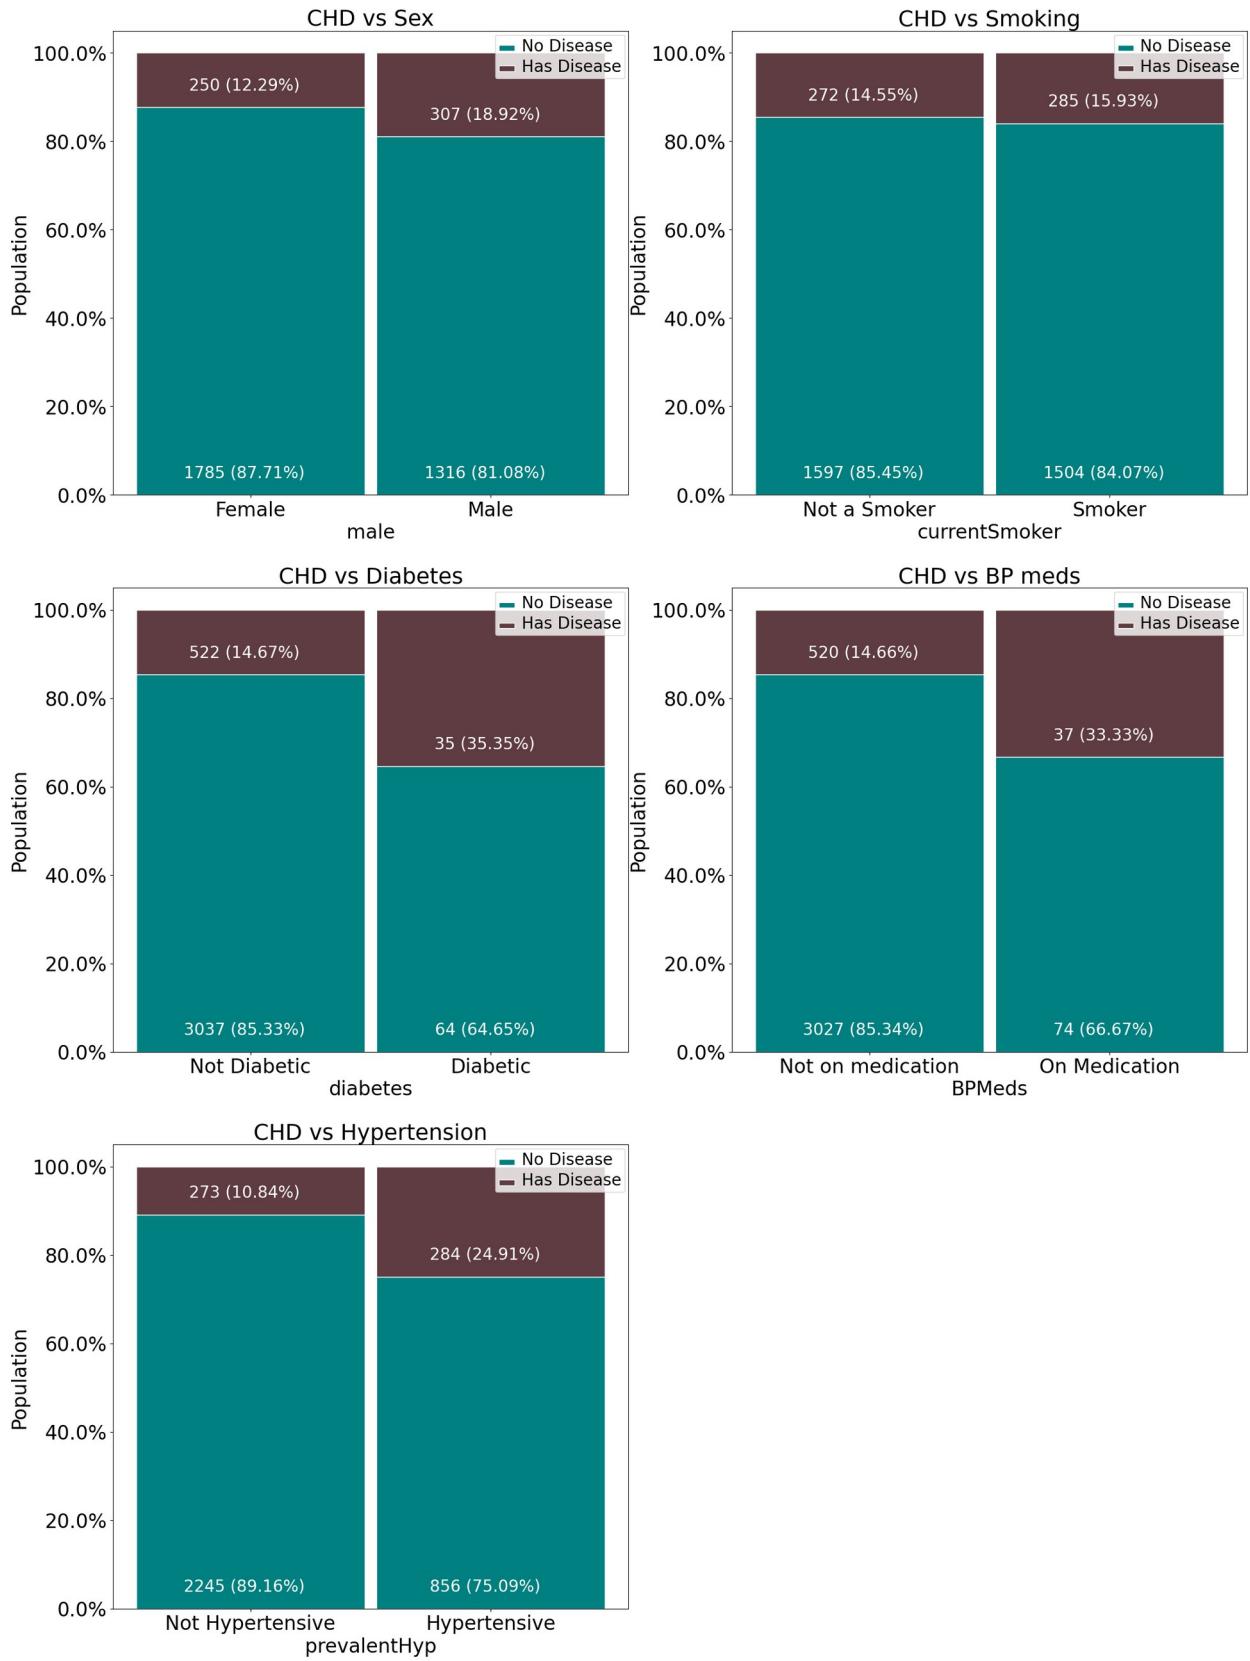
#draw BP meds vs disease outcome
plt.subplot(grid_rows, grid_cols, 4)
temp =
data[['BPMeds','TenYearCHD']].groupby(['BPMeds','TenYearCHD']).size()
.unstack('TenYearCHD')
temp.rename(index={0:'Not on medication', 1:'On Medication'},
columns={0:'No Disease', 1:'Has Disease'}, inplace = True)
stacked_barchart(temp, title = 'CHD vs BP meds', ylabel =
'Population')

#draw Hypertension vs disease outcome
plt.subplot(grid_rows, grid_cols, 5)
temp =
data[['prevalentHyp','TenYearCHD']].groupby(['prevalentHyp','TenYearCH
D']).size().unstack('TenYearCHD')
temp.rename(index={0:'Not Hypertensive', 1:'Hypertensive'},
columns={0:'No Disease', 1:'Has Disease'}, inplace = True)
stacked_barchart(temp, title = 'CHD vs Hypertension', ylabel =
'Population')

/tmp/ipython-input-2869839754.py:31: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.

```

```
    ax.set_yticklabels([str(y) + '%' for y in y_labels])
/tmp/ipython-input-2869839754.py:31: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
    ax.set_yticklabels([str(y) + '%' for y in y_labels])
/tmp/ipython-input-2869839754.py:31: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
    ax.set_yticklabels([str(y) + '%' for y in y_labels])
/tmp/ipython-input-2869839754.py:31: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
    ax.set_yticklabels([str(y) + '%' for y in y_labels])
/tmp/ipython-input-2869839754.py:31: UserWarning: set_ticklabels()
should only be used with a fixed number of ticks, i.e. after
set_ticks() or using a FixedLocator.
    ax.set_yticklabels([str(y) + '%' for y in y_labels])
```

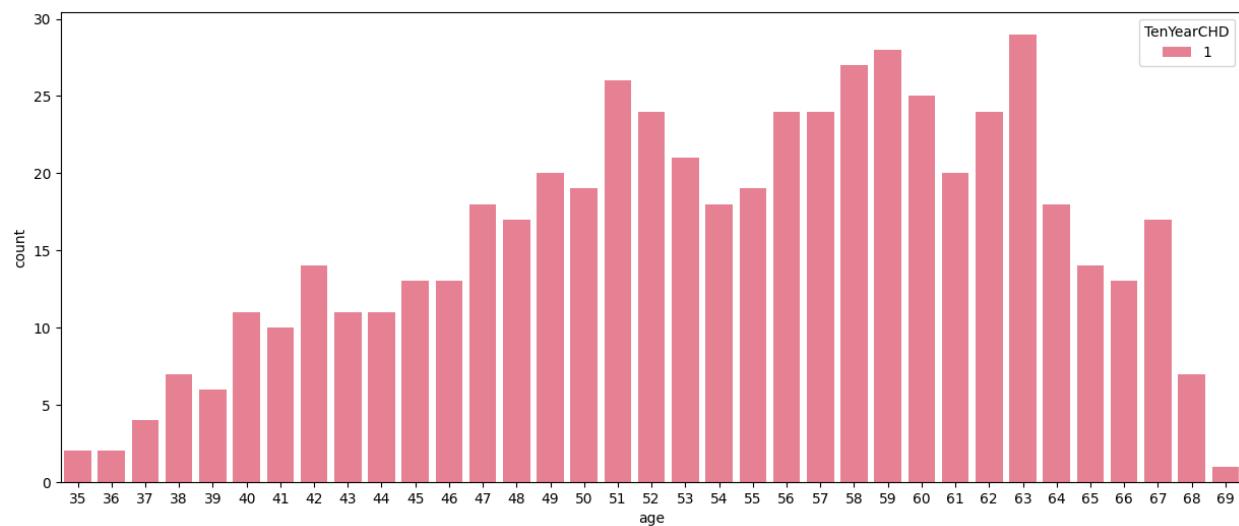


Due to the imbalanced nature of the dataset it is difficult to make conclusions but based on what is observed but these are the conclusions that can be drawn:

Slightly more males are suffering from CHD than females The percentage of people who have CHD is almost equal between smokers and non smokers The percentage of people who have CHD is higher among the diabetic, and those with prevalent hypertension as compared to those who dont have similar morbidities A larger percentage of the people who have CHD are on blood pressure medication

Number of people who have disease vs age

```
positive_cases = data[data['TenYearCHD'] == 1]
plt.figure(figsize=(15,6))
sns.countplot(x='age', data = positive_cases, hue = 'TenYearCHD',
palette='husl')
plt.show()
```

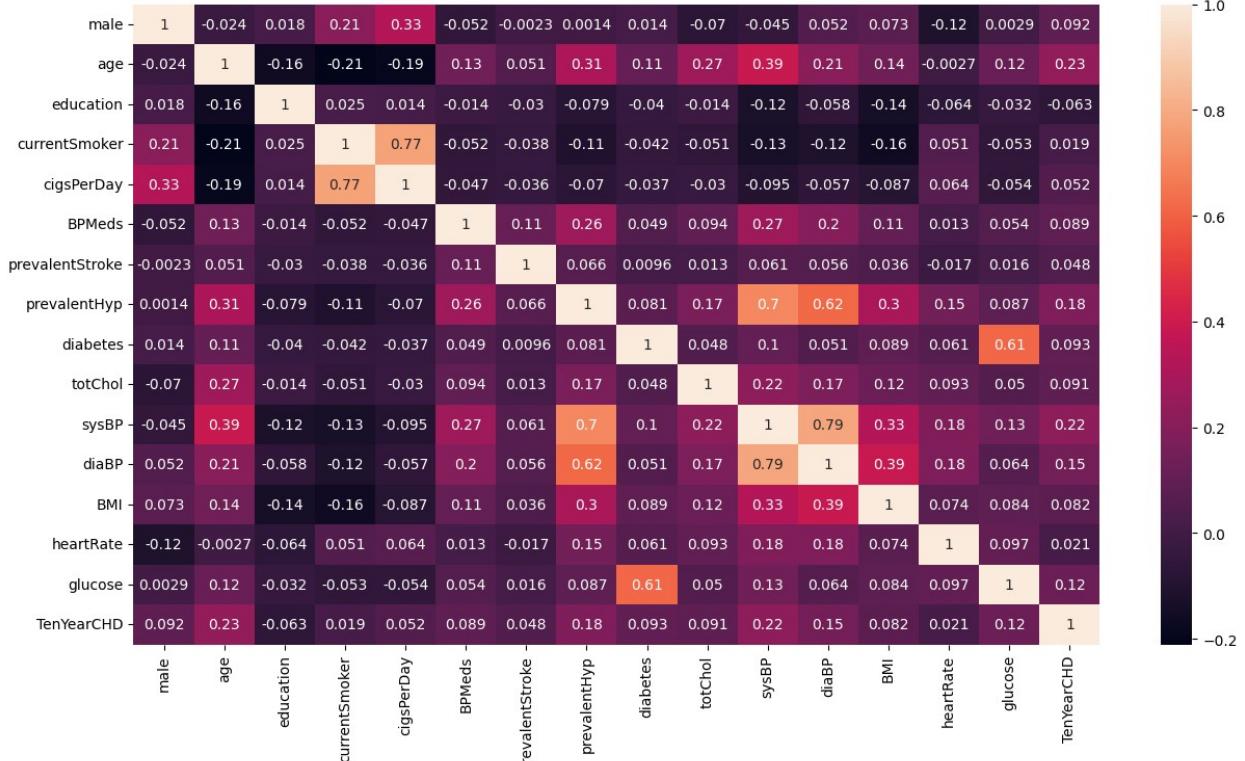


The people with the highest risk of developing CHD are betwwen the ages of 51 and 63

The number of sick people generally increases with age

Correlation Heat map

```
plt.figure(figsize=(15,8))
sns.heatmap(data.corr(), annot = True)
plt.show()
```



Feature Selection

```
from sklearn.ensemble import RandomForestClassifier
from boruta import BorutaPy

!pip install boruta

Requirement already satisfied: boruta in
/usr/local/lib/python3.12/dist-packages (0.4.3)
Requirement already satisfied: numpy>=1.10.4 in
/usr/local/lib/python3.12/dist-packages (from boruta) (2.0.2)
Requirement already satisfied: scikit-learn>=0.17.1 in
/usr/local/lib/python3.12/dist-packages (from boruta) (1.6.1)
Requirement already satisfied: scipy>=0.17.0 in
/usr/local/lib/python3.12/dist-packages (from boruta) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.17.1->boruta) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.17.1->boruta) (3.6.0)

#define the features
X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

```
forest = RandomForestClassifier(n_estimators=1000, n_jobs=-1,
class_weight='balanced')

# define Boruta feature selection method
feat_selector = BorutaPy(forest, n_estimators='auto', verbose=2)

# find all relevant features
feat_selector.fit(X, y)

Iteration:    1 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    2 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    3 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    4 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    5 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    6 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    7 / 100
Confirmed:    0
Tentative:    15
Rejected:    0
Iteration:    8 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    9 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    10 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    11 / 100
```

```
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 12 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 13 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 14 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 15 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 16 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 17 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 18 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 19 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 20 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 21 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 22 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 23 / 100
Confirmed: 2
```

```
Tentative: 4
Rejected: 9
Iteration: 24 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 25 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 26 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 27 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 28 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 29 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 30 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 31 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 32 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 33 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 34 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 35 / 100
Confirmed: 2
Tentative: 4
```

```
Rejected: 9
Iteration: 36 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 37 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 38 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 39 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 40 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 41 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 42 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 43 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 44 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 45 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 46 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 47 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
```

```
Iteration:    48 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    49 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    50 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    51 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    52 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    53 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    54 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    55 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    56 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    57 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    58 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    59 / 100
Confirmed:    2
Tentative:    4
Rejected:    9
Iteration:    60 / 100
```

```
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 61 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 62 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 63 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 64 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 65 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 66 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 67 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 68 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 69 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 70 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 71 / 100
Confirmed: 2
Tentative: 4
Rejected: 9
Iteration: 72 / 100
Confirmed: 2
```

```
Tentative: 3
Rejected: 10
Iteration: 73 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 74 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 75 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 76 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 77 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 78 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 79 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 80 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 81 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 82 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 83 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 84 / 100
Confirmed: 2
Tentative: 3
```

```
Rejected: 10
Iteration: 85 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 86 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 87 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 88 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 89 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 90 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 91 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 92 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 93 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 94 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 95 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
Iteration: 96 / 100
Confirmed: 2
Tentative: 3
Rejected: 10
```

```

Iteration:    97 / 100
Confirmed:    2
Tentative:    3
Rejected:    10
Iteration:    98 / 100
Confirmed:    2
Tentative:    3
Rejected:    10
Iteration:    99 / 100
Confirmed:    2
Tentative:    3
Rejected:    10

BorutaPy finished running.

Iteration:    100 / 100
Confirmed:    2
Tentative:    2
Rejected:    11

BorutaPy(estimator=RandomForestClassifier(class_weight='balanced',
                                             n_estimators=31, n_jobs=-1,
                                             random_state=RandomState(MT19937) at 0x7FD07CA80A40),
          n_estimators='auto',
          random_state=RandomState(MT19937) at 0x7FD07CA80A40,
          verbose=2)

# show the most important features
most_important = data.columns[:-1][feat_selector.support_].tolist()
most_important

['age', 'sysBP']

# select the top 6 features
top_features = data.columns[:-1][feat_selector.ranking_ <=6].tolist()
top_features

['age',
 'cigsPerDay',
 'totChol',
 'sysBP',
 'diaBP',
 'BMI',
 'heartRate',
 'glucose']

import statsmodels.api as sm
X_top = data[top_features]
y = data['TenYearCHD']

```

```

res = sm.Logit(y,X_top).fit()
res.summary()

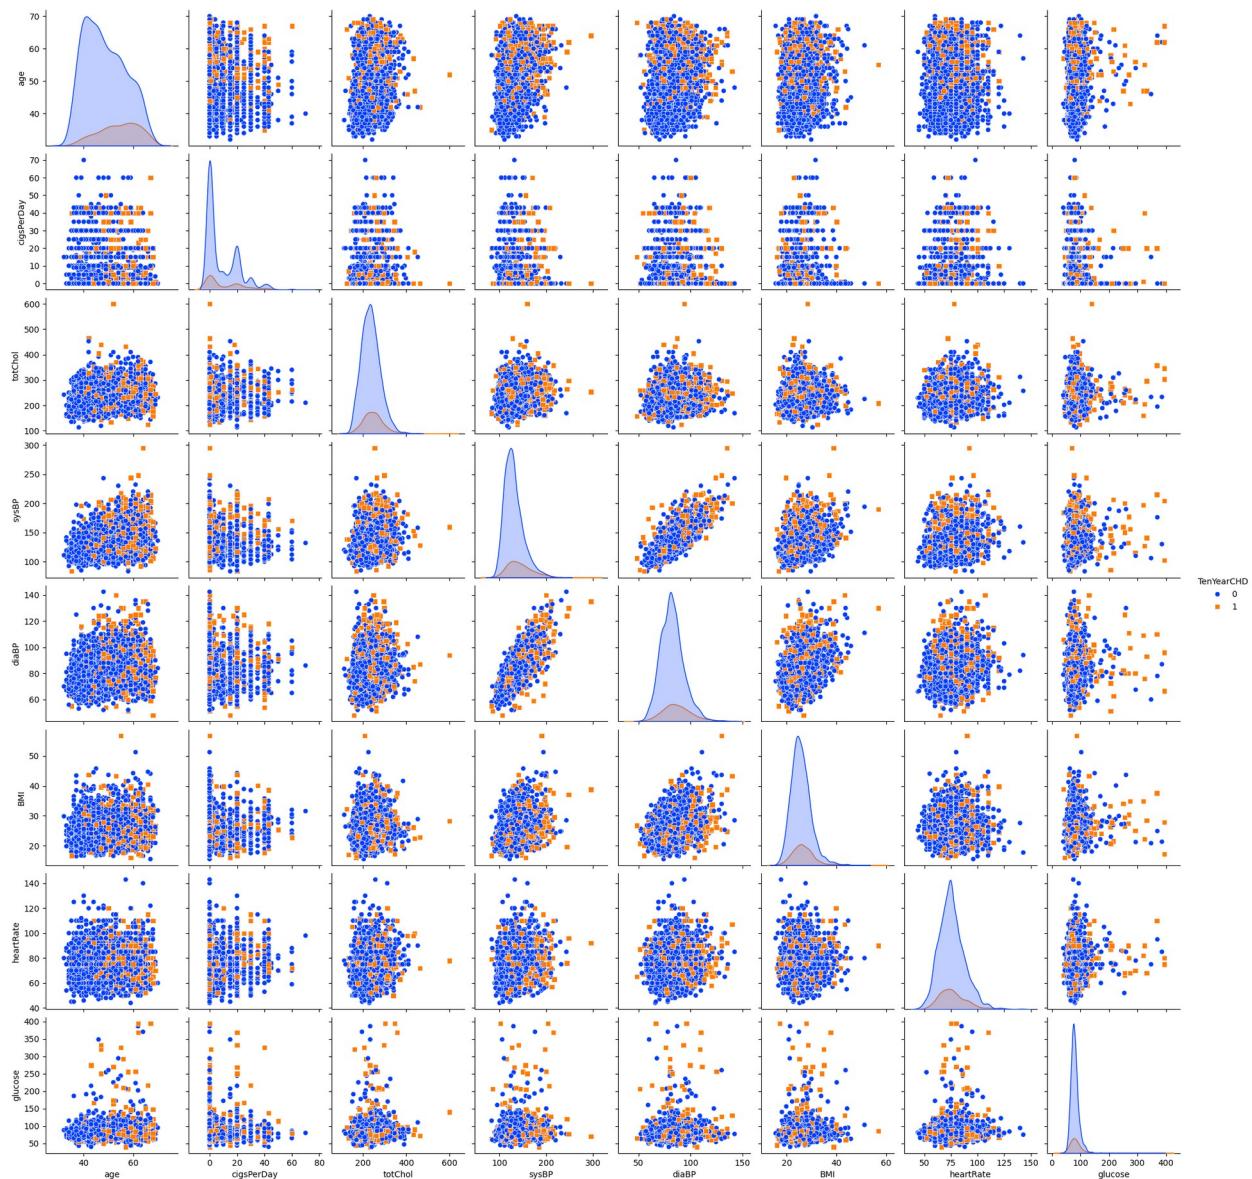
Optimization terminated successfully.
    Current function value: 0.413285
    Iterations 6

<class 'statsmodels.iolib.summary.Summary'>
"""
                                Logit Regression Results
=====
=====
Dep. Variable:          TenYearCHD    No. Observations:      3658
Model:                 Logit     Df Residuals:           3650
Method:                MLE      Df Model:             7
Date:      Tue, 28 Oct 2025   Pseudo R-squ.:  0.03127
Time:          16:41:15      Log-Likelihood: -1511.8
converged:            True     LL-Null:        -1560.6
Covariance Type:    nonrobust   LLR p-value: 3.388e-18
=====
=====

              coef    std err         z      P>|z|      [ 0.025
0.975]
-----
age          0.0259    0.006     4.516      0.000      0.015
0.037
cigsPerDay   0.0160    0.004     4.190      0.000      0.009
0.023
totChol     -0.0032    0.001    -2.907      0.004     -0.005
-0.001
sysBP        0.0262    0.003     7.761      0.000      0.020
0.033
diaBP       -0.0290    0.006    -4.838      0.000     -0.041
-0.017
BMI          -0.0517    0.012    -4.244      0.000     -0.076
-0.028
heartRate   -0.0333    0.004    -8.624      0.000     -0.041
-0.026
glucose      0.0041    0.002     2.470      0.014      0.001
0.007
=====
=====
```



```
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
/usr/local/lib/python3.12/dist-packages/seaborn/axisgrid.py:1615:  
UserWarning: The palette list has more values (10) than needed (2),  
which may not be intended.  
    func(x=x, y=y, **kwargs)  
  
<seaborn.axisgrid.PairGrid at 0x7fd03d1d2c60>
```



```

from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
from collections import Counter

X = data[top_features]
y = data.iloc[:, -1]

# the numbers before SMOTE
num_before = dict(Counter(y))

#perform SMOTE

# define pipeline
over = SMOTE(sampling_strategy=0.8)

```

```

under = RandomUnderSampler(sampling_strategy=0.8)
steps = [('o', over), ('u', under)]
pipeline = Pipeline(steps=steps)

# transform the dataset
X_smote, y_smote = pipeline.fit_resample(X, y)

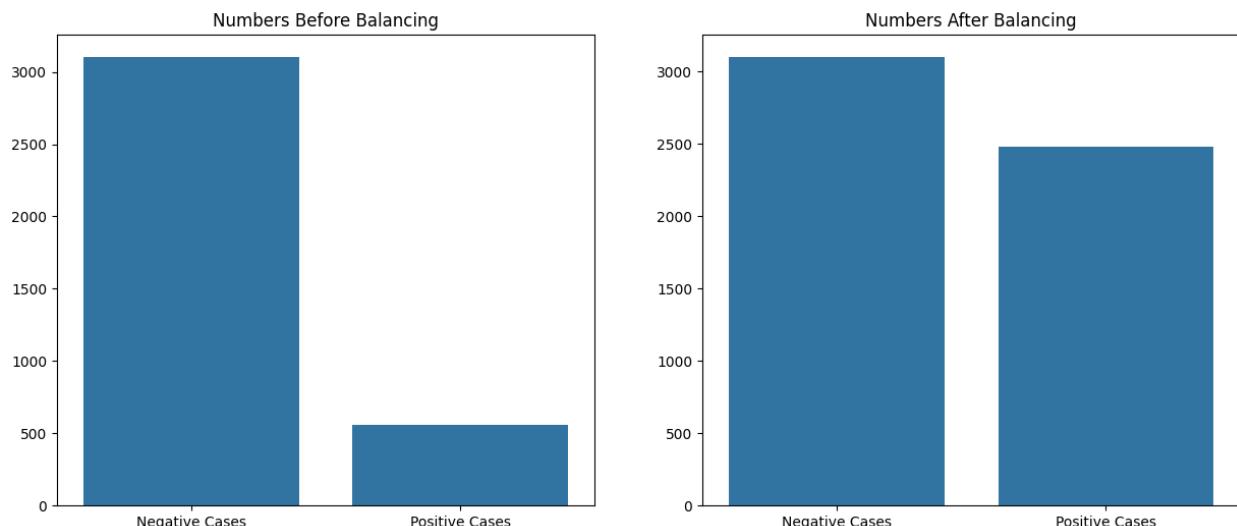
#the numbers after SMOTE
num_after =dict(Counter(y_smote))

print(num_before, num_after)

{0: 3101, 1: 557} {0: 3100, 1: 2480}

labels = ["Negative Cases","Positive Cases"]
plt.figure(figsize=(15,6))
plt.subplot(1,2,1)
sns.barplot(x=labels, y=list(num_before.values()))
plt.title("Numbers Before Balancing")
plt.subplot(1,2,2)
sns.barplot(x=labels, y=list(num_after.values()))
plt.title("Numbers After Balancing")
plt.show()

```



After applying SMOTE, the new dataset is much more balanced: the new ratio between negative and positive cases is 1:1.2 up from 1:5.57

Splitting data to Training and Testing set

```

# new dataset
new_data = pd.concat([pd.DataFrame(X_smote), pd.DataFrame(y_smote)],
axis=1)
new_data.columns = ['age', 'cigsPerDay', 'totChol', 'sysBP', 'diaBP',
'smoke', 'alco', 'active', 'stroke', 'hypertension', 'diabetes',
'heartDisease', 'cardio']

```

```
'BMI', 'heartRate', 'glucose', 'TenYearCHD']
new_data.head()

{"summary": {"name": "new_data", "rows": 5580, "fields": [{"column": "age", "dtype": "number", "min": 32, "max": 70, "std": 8, "num_unique_values": 39, "samples": [67, 66, 51], "semantic_type": "\\"}, {"column": "cigsPerDay", "dtype": "number", "min": 0.0, "max": 70.0, "std": 11.786497053317436, "num_unique_values": 1022, "samples": [13.565418159273838, 0.25182790036791936, 0.4359354099353465], "semantic_type": "\\"}, {"column": "totChol", "dtype": "number", "min": 113.0, "max": 600.0, "std": 44.999664220334225, "num_unique_values": 2078, "samples": [211.6772155023817, 182.31866094439306, 218.87492127878448], "semantic_type": "\\"}, {"column": "sysBP", "dtype": "number", "min": 83.5, "max": 295.0, "std": 23.54639666156191, "num_unique_values": 2111, "samples": [125.82533449921712, 113.71045859878971, 94.31280282047642], "semantic_type": "\\"}, {"column": "diaBP", "dtype": "number", "min": 48.0, "max": 142.5, "std": 12.439863788070305, "num_unique_values": 1994, "samples": [115.22328854038206, 77.89064890790087, 95.47487834044118], "semantic_type": "\\"}, {"column": "BMI", "dtype": "number", "min": 15.54, "max": 56.8, "std": 3.937757123899171, "num_unique_values": 3217, "samples": [31.35, 30.3, 30.591532451553633], "semantic_type": "\\"}, {"column": "heartRate", "dtype": "number", "min": 44.0, "max": 143.0, "std": 11.619510898872495, "num_unique_values": 1842, "samples": [86.59958491578152, 68.3990319413708, 78.57773073725667], "semantic_type": "\\"}, {"column": "glucose", "dtype": "number", "min": 40.0, "max": 394.0, "std": 29.258899439578457, "num_unique_values": 1971, "samples": [1971, 1971, 1971]}], "description": "\\"}, {"properties": [{"column": "age", "dtype": "number", "min": 32, "max": 70, "std": 8, "num_unique_values": 39, "samples": [67, 66, 51], "semantic_type": "\\"}, {"column": "cigsPerDay", "dtype": "number", "min": 0.0, "max": 70.0, "std": 11.786497053317436, "num_unique_values": 1022, "samples": [13.565418159273838, 0.25182790036791936, 0.4359354099353465], "semantic_type": "\\"}, {"column": "totChol", "dtype": "number", "min": 113.0, "max": 600.0, "std": 44.999664220334225, "num_unique_values": 2078, "samples": [211.6772155023817, 182.31866094439306, 218.87492127878448], "semantic_type": "\\"}, {"column": "sysBP", "dtype": "number", "min": 83.5, "max": 295.0, "std": 23.54639666156191, "num_unique_values": 2111, "samples": [125.82533449921712, 113.71045859878971, 94.31280282047642], "semantic_type": "\\"}, {"column": "diaBP", "dtype": "number", "min": 48.0, "max": 142.5, "std": 12.439863788070305, "num_unique_values": 1994, "samples": [115.22328854038206, 77.89064890790087, 95.47487834044118], "semantic_type": "\\"}, {"column": "BMI", "dtype": "number", "min": 15.54, "max": 56.8, "std": 3.937757123899171, "num_unique_values": 3217, "samples": [31.35, 30.3, 30.591532451553633], "semantic_type": "\\"}, {"column": "heartRate", "dtype": "number", "min": 44.0, "max": 143.0, "std": 11.619510898872495, "num_unique_values": 1842, "samples": [86.59958491578152, 68.3990319413708, 78.57773073725667], "semantic_type": "\\"}, {"column": "glucose", "dtype": "number", "min": 40.0, "max": 394.0, "std": 29.258899439578457, "num_unique_values": 1971, "samples": [1971, 1971, 1971]}], "description": "\\"}]}]
```

```

\"samples\": [\n    89.50118920631317,\n    73.64168764506331,\n    75.18603613470881\n],\n  ],\n  \"semantic_type\": \"\",\n  \"description\": \"\"\n},\n{\n  \"column\": \"TenYearCHD\",\n  \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 0,\n    \"min\": 0,\n    \"max\": 1,\n    \"num_unique_values\": 2,\n    \"samples\": [1,\n    0]\n  },\n  \"semantic_type\": \"\",\n  \"description\": \"\"\n}\n},\n\"type\":\"dataframe\",\"variable_name\":\"new_data\"}

X_new = new_data[top_features]
y_new= new_data.iloc[:, -1]
X_new.head()

{
  \"summary\": {\n    \"name\": \"X_new\", \"rows\": 5580,\n    \"fields\": [\n      {\n        \"column\": \"age\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 8,\n          \"min\": 32,\n          \"max\": 70,\n          \"num_unique_values\": 39,\n          \"samples\": [\n            67,\n            66,\n            51\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"cigsPerDay\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 11.786497053317436,\n          \"min\": 0.0,\n          \"max\": 70.0,\n          \"num_unique_values\": 1022,\n          \"samples\": [\n            13.565418159273838,\n            0.25182790036791936,\n            0.4359354099353465\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"totChol\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 44.999664220334225,\n          \"min\": 113.0,\n          \"max\": 600.0,\n          \"num_unique_values\": 2078,\n          \"samples\": [\n            182.31866094439306,\n            211.6772155023817,\n            218.87492127878448\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"sysBP\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 23.54639666156191,\n          \"min\": 83.5,\n          \"max\": 295.0,\n          \"num_unique_values\": 2111,\n          \"samples\": [\n            113.71045859878971,\n            125.82533449921712,\n            94.31280282047642\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"diaBP\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 12.439863788070305,\n          \"min\": 48.0,\n          \"max\": 142.5,\n          \"num_unique_values\": 1994,\n          \"samples\": [\n            77.89064890790087,\n            115.22328854038206,\n            95.47487834044118\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"BMI\", \"properties\": {\n          \"dtype\": \"number\", \"std\": 3.937757123899171,\n          \"min\": 15.54,\n          \"max\": 56.8,\n          \"num_unique_values\": 3217,\n          \"samples\": [\n            1,\n            0\n          ],\n          \"semantic_type\": \"\", \"description\": \"\"\n        }\n      }\n    ],\n    \"type\": \"dataframe\", \"variable_name\": \"new_data\"\n  }
}

```

```

30.3,\n            30.591532451553633,\n            31.35\n        ],\n    {\\"semantic_type\\": \"\",\\n        \\"description\\": \"\\n        \"},\\n    {\\"column\\": \"heartRate\",\\n        \\"properties\\\": {\n            \\"dtype\\": \"number\",\\n            \\"std\\": 11.619510898872495,\n            \\"min\\": 44.0,\n            \\"max\\": 143.0,\n            \\"num_unique_values\\": 1842,\n            \\"samples\\\": [\n                68.3990319413708,\n                86.59958491578152,\n                78.57773073725667\n            ],\n            \\"semantic_type\\": \"\",\\n            \\"description\\": \"\\n            \"},\\n            {\n                \\"column\\": \"glucose\",\\n                \\"properties\\\": {\n                    \\"dtype\\": \"number\",\\n                    \\"std\\": 29.258899439578457,\n                    \\"min\\": 40.0,\n                    \\"max\\": 394.0,\n                    \\"num_unique_values\\": 1971,\n                    \\"samples\\\": [\n                        89.50118920631317,\n                        73.64168764506331,\n                        75.18603613470881\n                    ],\n                    \\"semantic_type\\": \"\",\\n                    \\"description\\": \"\\n                    \"},\\n                }\n            ]\\n        },\\n        \\"type\\": \"dataframe\",\\n        \\"variable_name\\": \"X_new\"\n    }\n\n# split the dataset\nfrom sklearn.model_selection import train_test_split\nX_train,X_test,y_train,y_test =\ntrain_test_split(X_new,y_new,test_size=.2,random_state=42)\n\nfrom sklearn.preprocessing import StandardScaler\nscaler = StandardScaler()\n\nX_train_scaled = scaler.fit_transform(X_train)\nX_train = pd.DataFrame(X_train_scaled)\n\nX_test_scaled = scaler.transform(X_test)\nX_test = pd.DataFrame(X_test_scaled)

```

Logistic regression

```

from sklearn.linear_model import LogisticRegression\nfrom sklearn.model_selection import GridSearchCV\nfrom sklearn.metrics import confusion_matrix\nfrom sklearn.metrics import accuracy_score\nfrom sklearn.metrics import f1_score\nfrom sklearn.metrics import classification_report\nfrom sklearn.metrics import\nrecall_score,precision_score,classification_report,roc_auc_score,roc_c\nurve\n\n# search for optimum parameters using gridsearch\nparams = {'penalty':['l1','l2'],\n          'C':[0.01,0.1,1,10,100],\n          'class_weight':[\"balanced\",None]}\nlogistic_clf =\nGridSearchCV(LogisticRegression(),param_grid=params,cv=10)

```

```
#train the classifier
logistic_clf.fit(X_train,y_train)

logistic_clf.best_params_

/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/
_validation.py:528: FitFailedWarning:
100 fits failed out of a total of 200.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

Below are more details about the failures:
-----
-----
100 fits failed with the following error:
Traceback (most recent call last):
  File
"/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_vali
dation.py", line 866, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.12/dist-packages/sklearn/base.py", line
1389, in wrapper
    return fit_method(estimator, *args, **kwargs)
          ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

  File
"/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logisti
c.py", line 1193, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
          ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

  File
"/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logisti
c.py", line 63, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1
penalty.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_searc
h.py:1108: UserWarning: One or more of the test scores are non-finite:
[      nan  0.66218387      nan  0.65367021      nan  0.661286
      nan  0.65569015      nan  0.66083858      nan  0.65658701
      nan  0.66106279      nan  0.65658701      nan  0.66106279
      nan  0.65658701]
warnings.warn(
{'C': 0.01, 'class_weight': 'balanced', 'penalty': 'l2'}
```

```

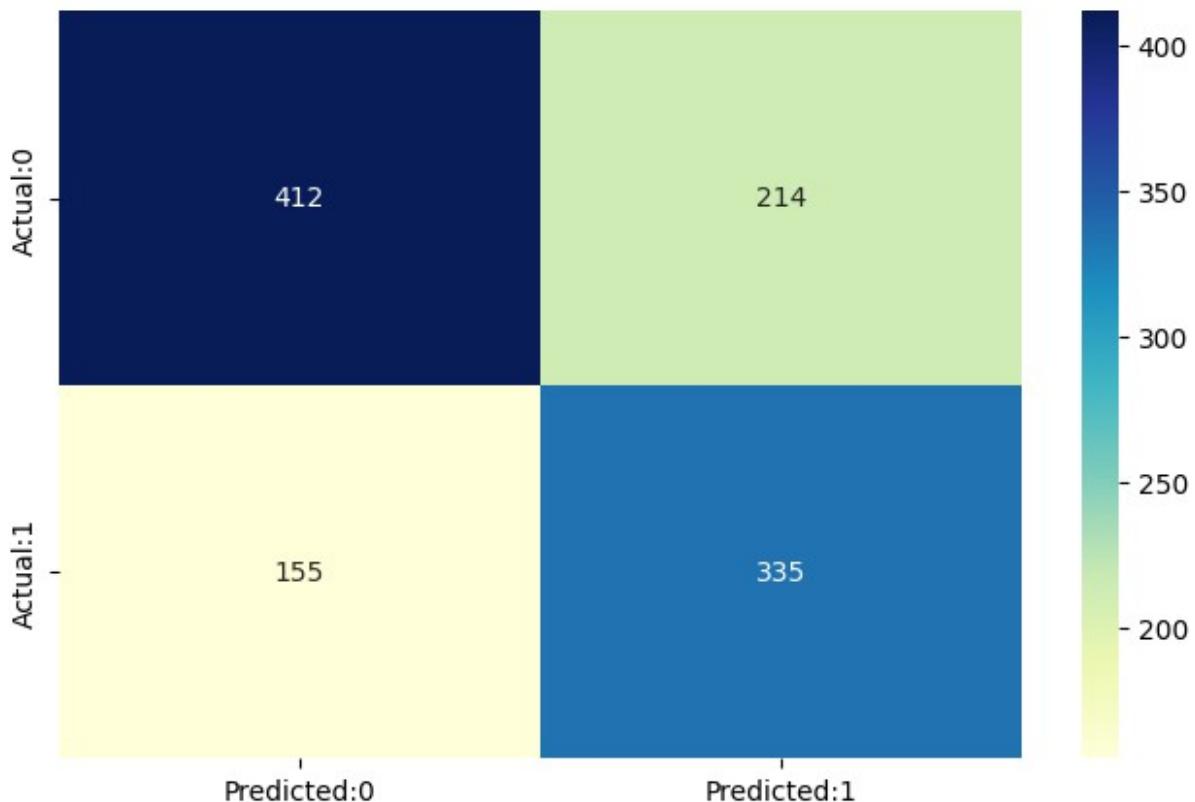
#make predictions
logistic_predict = logistic_clf.predict(X_test)
log_accuracy = accuracy_score(y_test,logistic_predict)
print(f"Using logistic regression we get an accuracy of {round(log_accuracy*100,2)}%")

Using logistic regression we get an accuracy of 66.94%

cm=confusion_matrix(y_test,logistic_predict)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],
 ,index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")

<Axes: >

```



```
print(classification_report(y_test,logistic_predict))
```

	precision	recall	f1-score	support
0	0.73	0.66	0.69	626
1	0.61	0.68	0.64	490
accuracy			0.67	1116
macro avg	0.67	0.67	0.67	1116

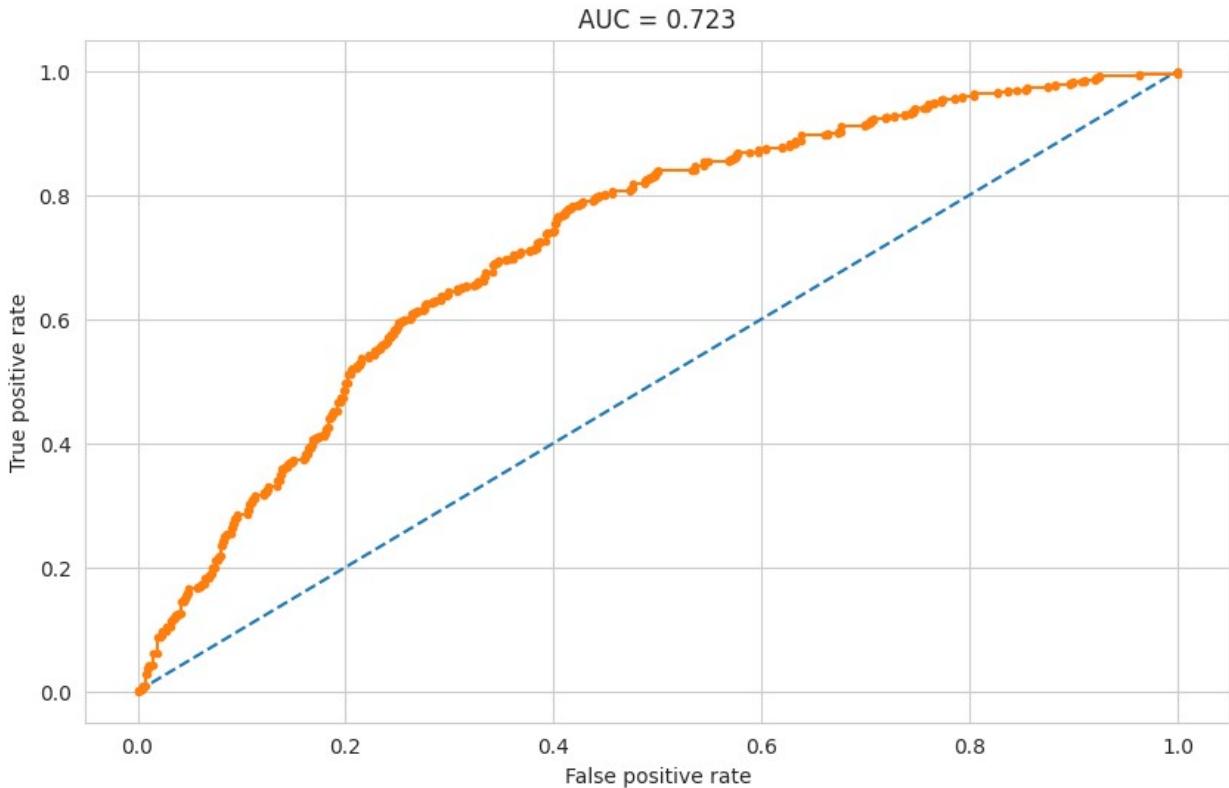
weighted avg	0.68	0.67	0.67	1116
--------------	------	------	------	------

```
logistic_f1 = f1_score(y_test, logistic_predict)
print(f'The f1 score for logistic regression is
{round(logistic_f1*100,2)}%)
```

The f1 score for logistic regression is 64.49%

```
# ROC curve and AUC
probs = logistic_clf.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
log_auc = roc_auc_score(y_test, probs)

# calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, probs)
# plot curve
sns.set_style('whitegrid')
plt.figure(figsize=(10,6))
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.ylabel('True positive rate')
plt.xlabel('False positive rate')
plt.title(f"AUC = {round(log_auc,3)}")
plt.show()
```

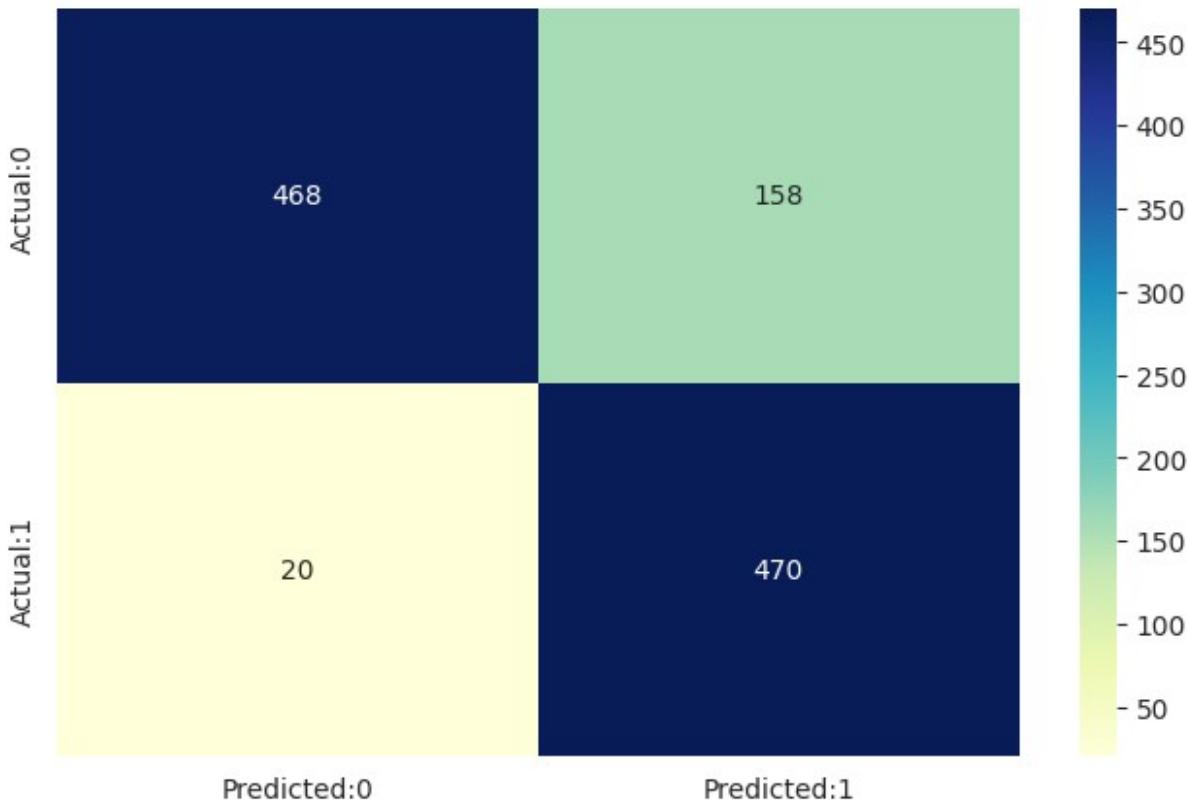


k-Nearest Neighbours

```
from sklearn.neighbors import KNeighborsClassifier
# search for optimum parameters using gridsearch
params= {'n_neighbors': np.arange(1, 10)}
grid_search = GridSearchCV(estimator = KNeighborsClassifier(),
param_grid = params,
scoring = 'accuracy', cv = 10, n_jobs = -1)
knn_clf = GridSearchCV(KNeighborsClassifier(),params,cv=3, n_jobs=-1)
# train the model
knn_clf.fit(X_train,y_train)
knn_clf.best_params_
# predictions
knn_predict = knn_clf.predict(X_test)
#accuracy
knn_accuracy = accuracy_score(y_test,knn_predict)
print(f"Using k-nearest neighbours we get an accuracy of
{round(knn_accuracy*100,2)}%")
cm=confusion_matrix(y_test,knn_predict)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1']
,index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
```

Using k-nearest neighbours we get an accuracy of 84.05%

```
<Axes: >
```



```
print(classification_report(y_test,knn_predict))
```

	precision	recall	f1-score	support
0	0.96	0.75	0.84	626
1	0.75	0.96	0.84	490
accuracy			0.84	1116
macro avg	0.85	0.85	0.84	1116
weighted avg	0.87	0.84	0.84	1116

```
knn_f1 = f1_score(y_test, knn_predict)
print(f'The f1 score for K nearest neighbours is {round(knn_f1*100,2)} %')
```

```
The f1 score for K nearest neighbours is 84.08%
```

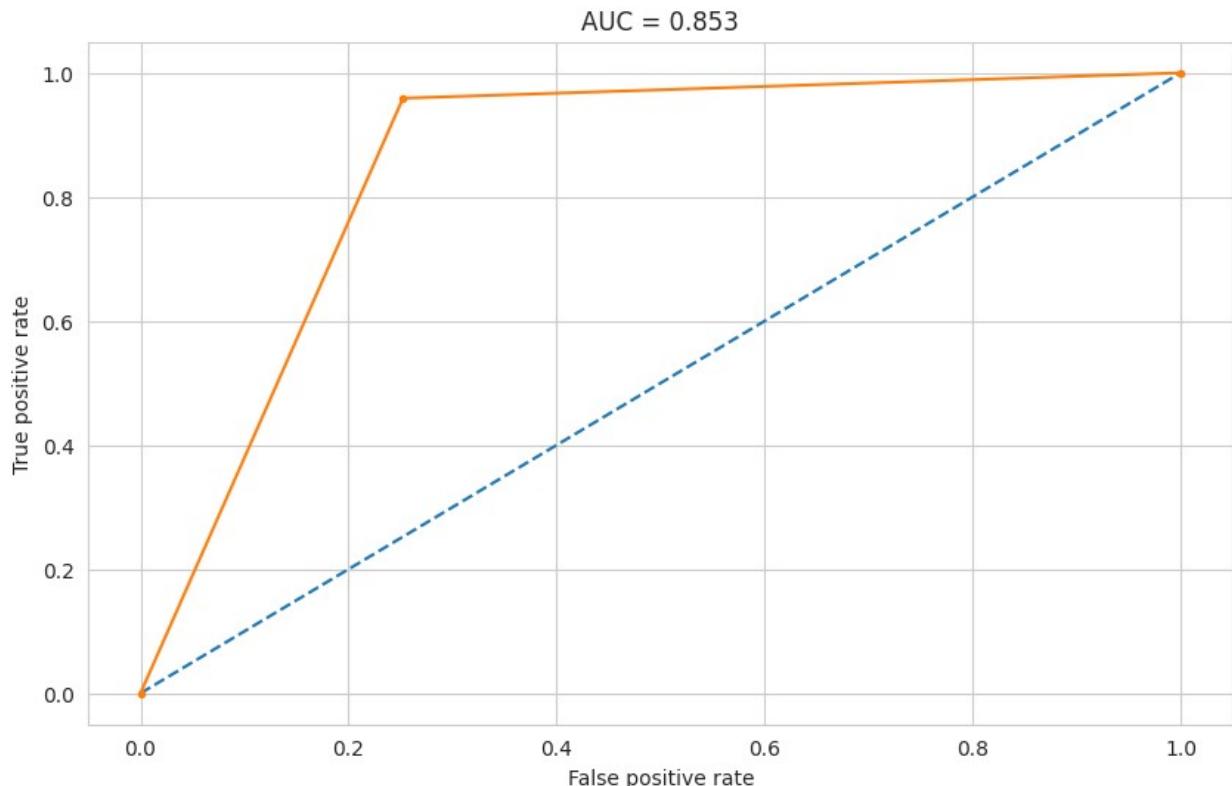
```
# ROC curve and AUC
probs = knn_clf.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
```

```

knn_auc = roc_auc_score(y_test, probs)

# calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, probs)
# plot curve
sns.set_style('whitegrid')
plt.figure(figsize=(10,6))
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.ylabel('True positive rate')
plt.xlabel('False positive rate')
plt.title(f"AUC = {round(knn_auc,3)}")
plt.show()

```



Decision Trees

```

from sklearn.tree import DecisionTreeClassifier
dtree= DecisionTreeClassifier(random_state=7)
# grid search for optimum parameters
params = {'max_features': ['auto', 'sqrt', 'log2'],
          'min_samples_split': [2,3,4,5,6,7,8,9,10,11,12,13,14,15],
          'min_samples_leaf':[1,2,3,4,5,6,7,8,9,10,11]}
tree_clf = GridSearchCV(dtree, param_grid=params, n_jobs=-1)
# train the model
tree_clf.fit(X_train,y_train)

```

```

tree_clf.best_params_
# predictions
tree_predict = tree_clf.predict(X_test)
#accuracy
tree_accuracy = accuracy_score(y_test,tree_predict)
print(f"Using Decision Trees we get an accuracy of
{round(tree_accuracy*100,2)}%")

Using Decision Trees we get an accuracy of 75.09%

/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/
_validation.py:528: FitFailedWarning:
770 fits failed out of a total of 2310.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.
```

Below are more details about the failures:

770 fits failed with the following error:
Traceback (most recent call last):
 File
 "/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_vali
 dation.py", line 866, in _fit_and_score
 estimator.fit(X_train, y_train, **fit_params)
 File "/usr/local/lib/python3.12/dist-packages/sklearn/base.py", line
1382, in wrapper
 estimator._validate_params()
 File "/usr/local/lib/python3.12/dist-packages/sklearn/base.py", line
436, in _validate_params
 validate_parameter_constraints()
 File
 "/usr/local/lib/python3.12/dist-packages/sklearn/utils/_param_validati
 on.py", line 98, in validate_parameter_constraints
 raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of DecisionTreeClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

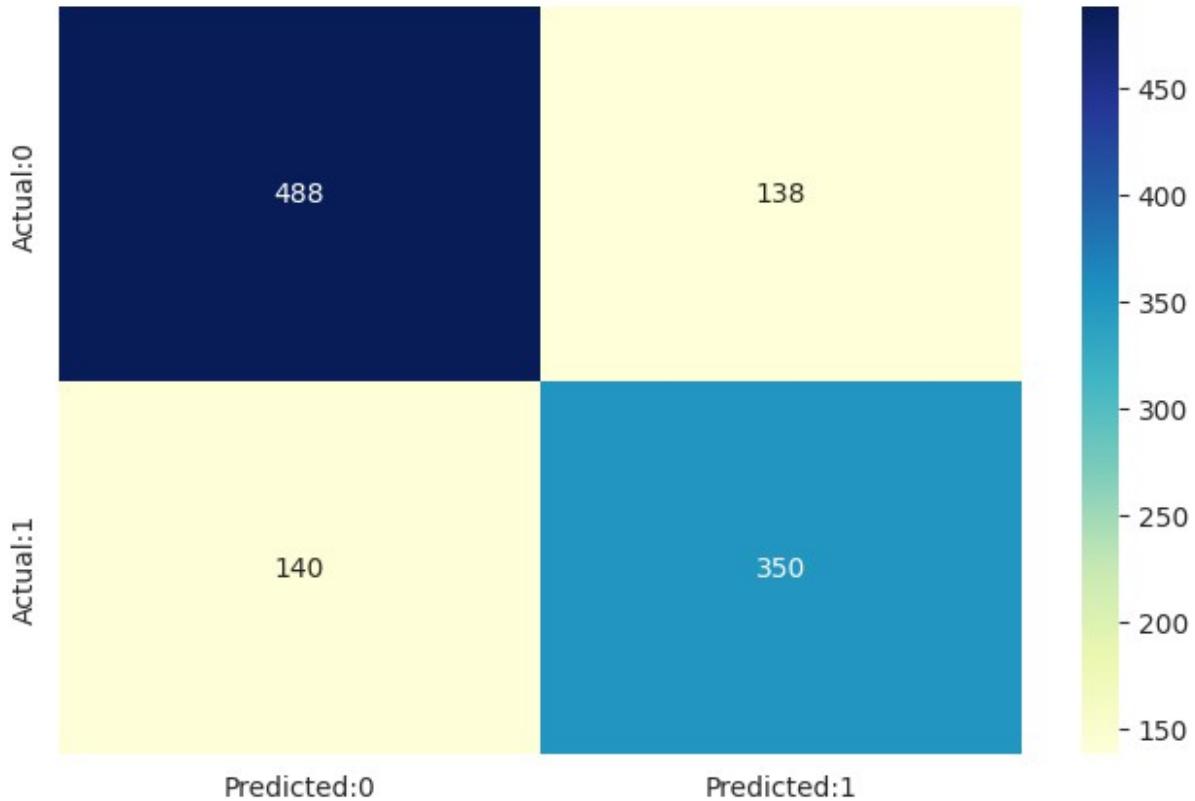
 warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.12/dist-packages/sklearn/model_selection/_searc
h.py:1108: UserWarning: One or more of the test scores are non-finite:
[nan nan nan nan nan nan
 nan nan nan nan nan nan

nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	0.73095978	0.73073431
0.71998855	0.72423985	0.71371529	0.71057553	0.7146109	0.7177459
0.71819433	0.7172947	0.71102697	0.72760559	0.71303712	0.70878206
0.72401413	0.72401413	0.72401413	0.71685104	0.71953158	0.72782227
0.71684552	0.70407755	0.70811067	0.71438794	0.71550199	0.71841452
0.70900426	0.7083384	0.72087487	0.72087487	0.72087487	0.72087487
0.72087487	0.71952857	0.71057452	0.7063177	0.70250805	0.71864703
0.71303637	0.71034981	0.70811368	0.71886245	0.71617463	0.71617463
0.71617463	0.71617463	0.71617463	0.71617463	0.71617463	0.71101743
0.71124717	0.71147314	0.69802851	0.70564606	0.71303135	0.71326109
0.7038566	0.7038566	0.7038566	0.7038566	0.7038566	0.7038566
0.7038566	0.7038566	0.7038566	0.70273302	0.7119173	0.7002674
0.70228609	0.71258669	0.70609223	0.70609223	0.70609223	0.70609223
0.70609223	0.70609223	0.70609223	0.70609223	0.70609223	0.70609223
0.70609223	0.7022871	0.69578787	0.69355224	0.69959752	0.69959752
0.69959752	0.69959752	0.69959752	0.69959752	0.69959752	0.69959752
0.69959752	0.69959752	0.69959752	0.69959752	0.69959752	0.69354697
0.69354873	0.69354873	0.69354873	0.69354873	0.69354873	0.69354873
0.69354873	0.69354873	0.69354873	0.69354873	0.69354873	0.69354873
0.69354873	0.69354873	0.67965843	0.67965843	0.67965843	0.67965843
0.67965843	0.67965843	0.67965843	0.67965843	0.67965843	0.67965843
0.67965843	0.67965843	0.67965843	0.67965843	0.68839554	0.68839554
0.68839554	0.68839554	0.68839554	0.68839554	0.68839554	0.68839554
0.68839554	0.68839554	0.68839554	0.68839554	0.68839554	0.68839554
0.69018751	0.69018751	0.69018751	0.69018751	0.69018751	0.69018751
0.69018751	0.69018751	0.69018751	0.69018751	0.69018751	0.69018751
0.69018751	0.69018751	0.72356972	0.73745324	0.7235647	0.72088089
0.73611196	0.71595318	0.71730374	0.71572695	0.7150533	0.71975781
0.71079673	0.71797112	0.71079999	0.70453201	0.72670547	0.72670547

```

0.72670547 0.71729219 0.71998478 0.71528078 0.72178026 0.72020649
0.71371002 0.70564932 0.71303964 0.70474894 0.71192107 0.71394001
0.71438392 0.71438392 0.71438392 0.71438392 0.71438392 0.72020975
0.71729696 0.72311752 0.72557862 0.71729345 0.71595243 0.71236724
0.71707526 0.71281442 0.69870342 0.69870342 0.69870342 0.69870342
0.69870342 0.69870342 0.69870342 0.71863422 0.70564555 0.70945671
0.71259346 0.72088089 0.70743852 0.71953384 0.7211026 0.7211026
0.7211026 0.7211026 0.7211026 0.7211026 0.7211026 0.7211026
0.7211026 0.70094231 0.70855885 0.7148316 0.707665 0.70900351
0.70878381 0.70878381 0.70878381 0.70878381 0.70878381 0.70878381
0.70878381 0.70878381 0.70878381 0.70878381 0.70878381 0.70833513
0.70788771 0.7038551 0.71505155 0.71505155 0.71505155 0.71505155
0.71505155 0.71505155 0.71505155 0.71505155 0.71505155 0.71505155
0.71505155 0.71505155 0.71505155 0.69915612 0.70094256 0.70094256
0.70094256 0.70094256 0.70094256 0.70094256 0.70094256 0.70094256
0.70094256 0.70094256 0.70094256 0.70094256 0.70094256 0.70094256
0.68840333 0.68840333 0.68840333 0.68840333 0.68840333 0.68840333
0.68840333 0.68840333 0.68840333 0.68840333 0.68840333 0.68840333
0.68840333 0.68840333 0.69735838 0.69735838 0.69735838 0.69735838
0.69735838 0.69735838 0.69735838 0.69735838 0.69735838 0.69735838
0.69735838 0.69735838 0.69735838 0.69735838 0.68458891 0.68458891
0.68458891 0.68458891 0.68458891 0.68458891 0.68458891 0.68458891
0.68458891 0.68458891 0.68458891 0.68458891 0.68458891 0.68458891]
warnings.warn(
cm=confusion_matrix(y_test,tree_predict)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1']
, index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")
<Axes: >

```



```

print(classification_report(y_test,tree_predict))

      precision    recall  f1-score   support

          0       0.78      0.78      0.78      626
          1       0.72      0.71      0.72      490

   accuracy                           0.75      1116
  macro avg       0.75      0.75      0.75      1116
weighted avg       0.75      0.75      0.75      1116

tree_f1 = f1_score(y_test, tree_predict)
print(f'The f1 score Descision trees is {round(tree_f1*100,2)}%')

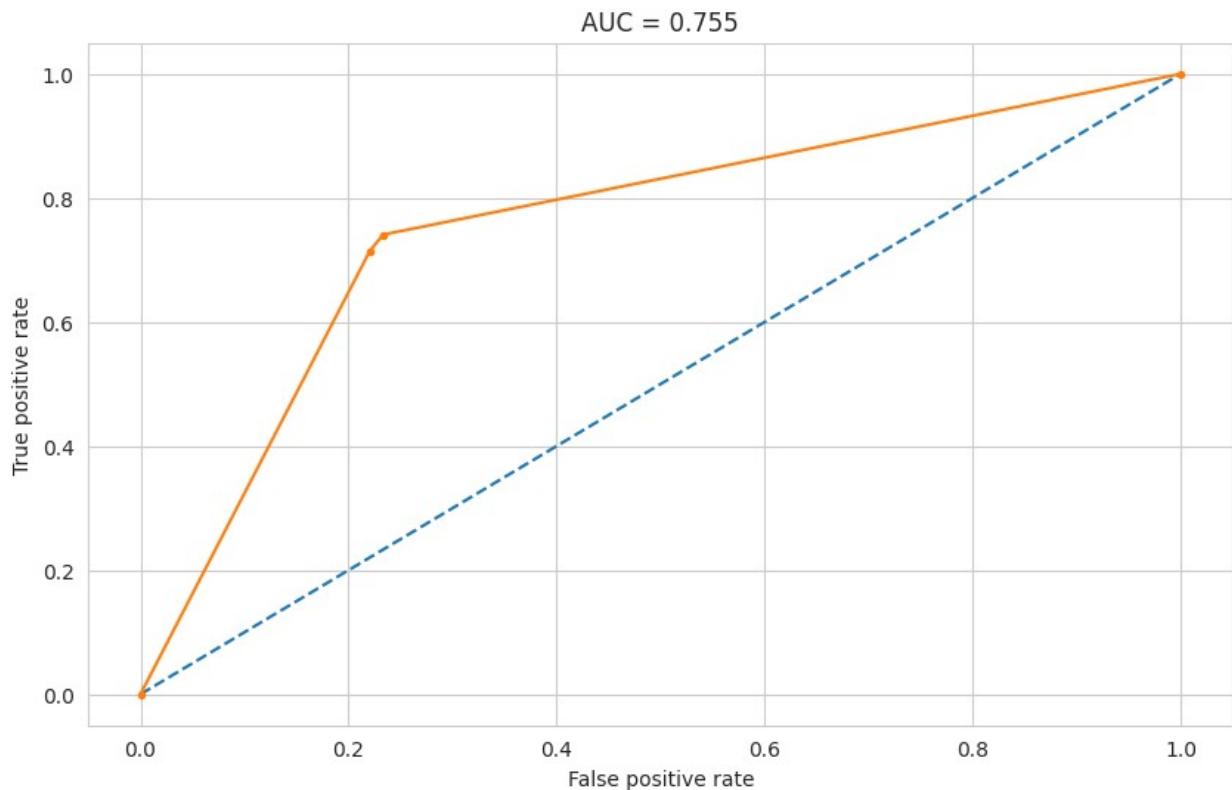
The f1 score Descision trees is 71.57%

# ROC curve and AUC
probs = tree_clf.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
tree_auc = roc_auc_score(y_test, probs)

# calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, probs)

```

```
# plot curve
sns.set_style('whitegrid')
plt.figure(figsize=(10,6))
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.ylabel('True positive rate')
plt.xlabel('False positive rate')
plt.title(f"AUC = {round(tree_auc,3)}")
plt.show()
```



Support Vector Machine

```
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV

#grid search for optimum parameters
Cs = [0.001, 0.01, 0.1, 1, 10]
gammas = [0.001, 0.01, 0.1, 1]
param_grid = {'C': Cs, 'gamma' : gammas}
svm_clf = GridSearchCV(SVC(kernel='rbf', probability=True),
param_grid, cv=10)

# train the model
svm_clf.fit(X_train,y_train)
svm_clf.best_params_
```

```

{'C': 10, 'gamma': 1}

# predictions
svm_predict = svm_clf.predict(X_test)

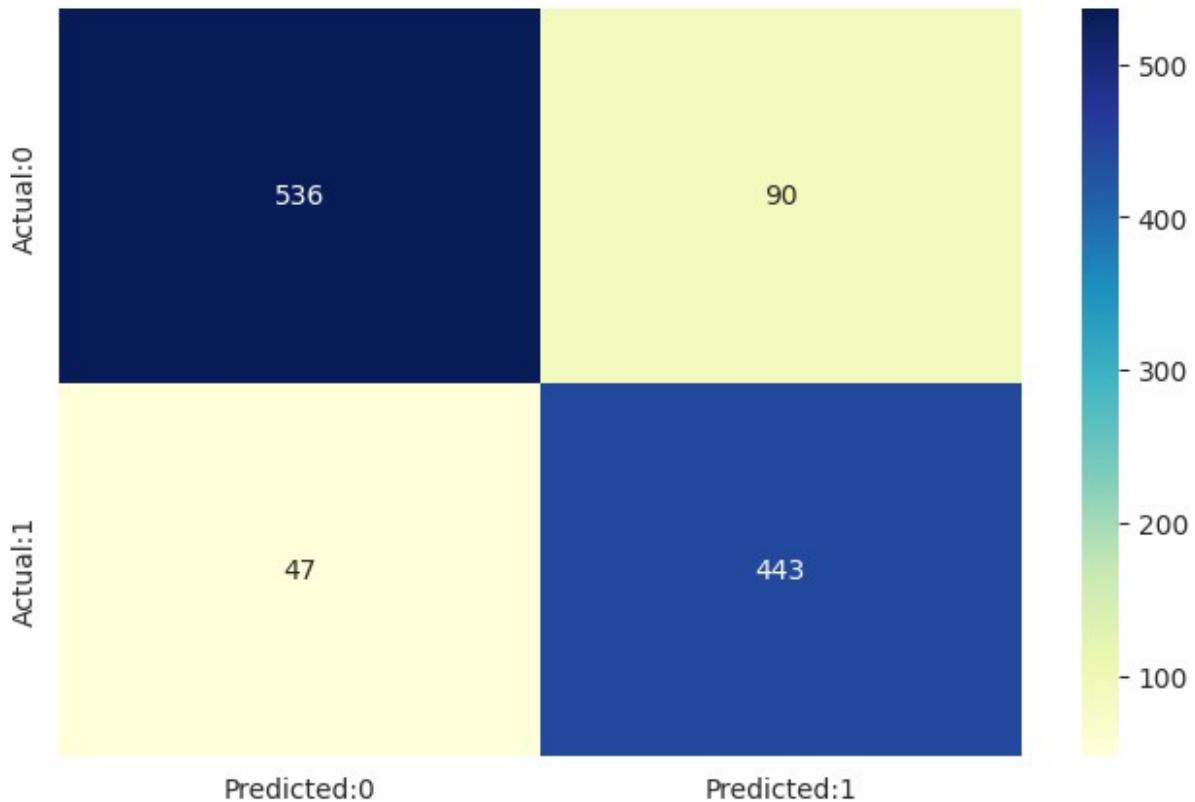
#accuracy
svm_accuracy = accuracy_score(y_test,svm_predict)
print(f"Using SVM we get an accuracy of {round(svm_accuracy*100,2)}%")

Using SVM we get an accuracy of 87.72%

cm=confusion_matrix(y_test,svm_predict)
conf_matrix=pd.DataFrame(data=cm,columns=['Predicted:0','Predicted:1'],
 ,index=['Actual:0','Actual:1'])
plt.figure(figsize = (8,5))
sns.heatmap(conf_matrix, annot=True,fmt='d',cmap="YlGnBu")

<Axes: >

```



```

print(classification_report(y_test,svm_predict))

      precision    recall  f1-score   support

       0       0.92      0.86      0.89      626
       1       0.83      0.90      0.87      490

```

	accuracy	0.88	0.88	0.88	1116
macro avg	0.88	0.88	0.88	1116	
weighted avg	0.88	0.88	0.88	1116	

```

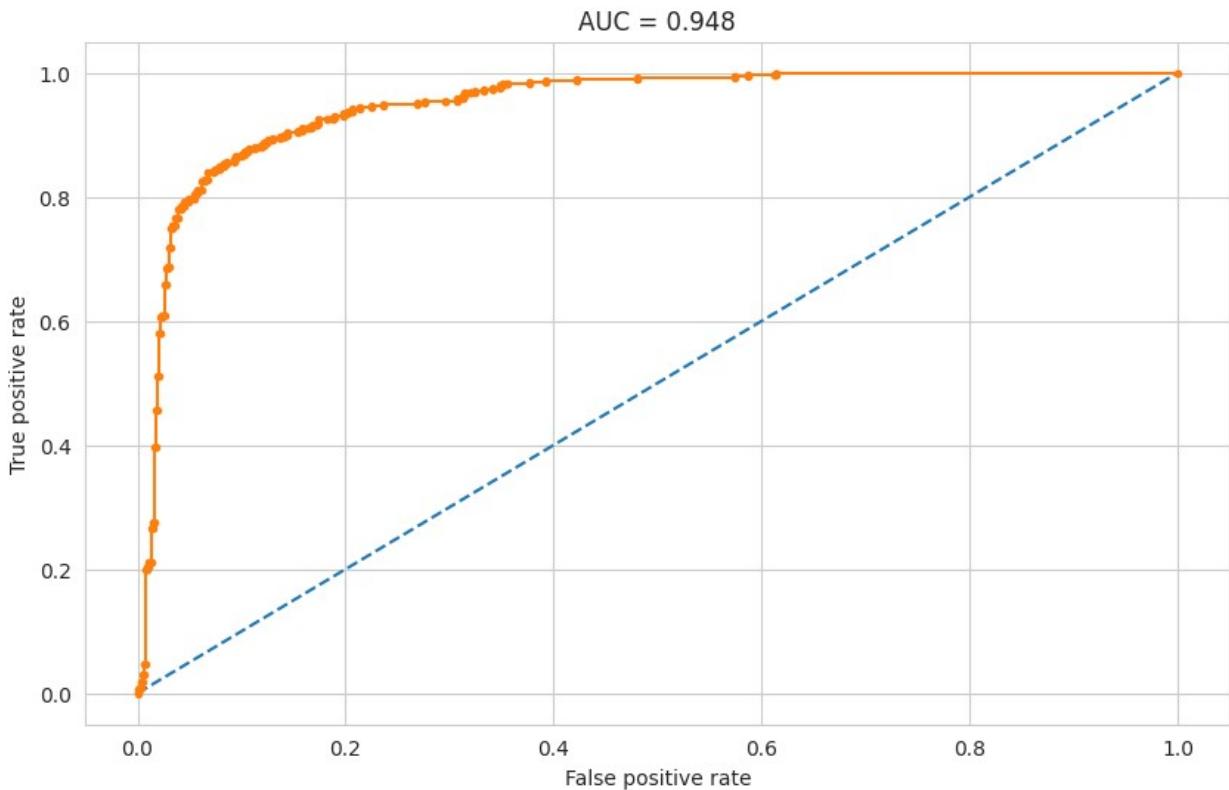
svm_f1 = f1_score(y_test, svm_predict)
print(f'The f1 score for SVM is {round(svm_f1*100,2)}%')

The f1 score for SVM is 86.61%

# ROC curve and AUC
probs = svm_clf.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
svm_auc = roc_auc_score(y_test, probs)

# calculate roc curve
fpr, tpr, thresholds = roc_curve(y_test, probs)
# plot curve
sns.set_style('whitegrid')
plt.figure(figsize=(10,6))
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.ylabel('True positive rate')
plt.xlabel('False positive rate')
plt.title(f"AUC = {round(svm_auc,3)}")
plt.show()

```



Random Forest

```

from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb

# --- Random Forest Classifier ---
rf_clf = RandomForestClassifier(
    n_estimators=200,
    max_depth=10,
    class_weight='balanced',
    random_state=42
)
rf_clf.fit(X_train, y_train)
rf_pred = rf_clf.predict(X_test)

rf_accuracy = accuracy_score(y_test, rf_pred)
rf_f1 = f1_score(y_test, rf_pred)
print(f"Using Random Forest we get an accuracy of {round(rf_accuracy*100,2)}%")
print(f"The F1 score for Random Forest is {round(rf_f1*100,2)}%")

Using Random Forest we get an accuracy of 77.78%
The F1 score for Random Forest is 76.87%

# Confusion Matrix
cm = confusion_matrix(y_test, rf_pred)

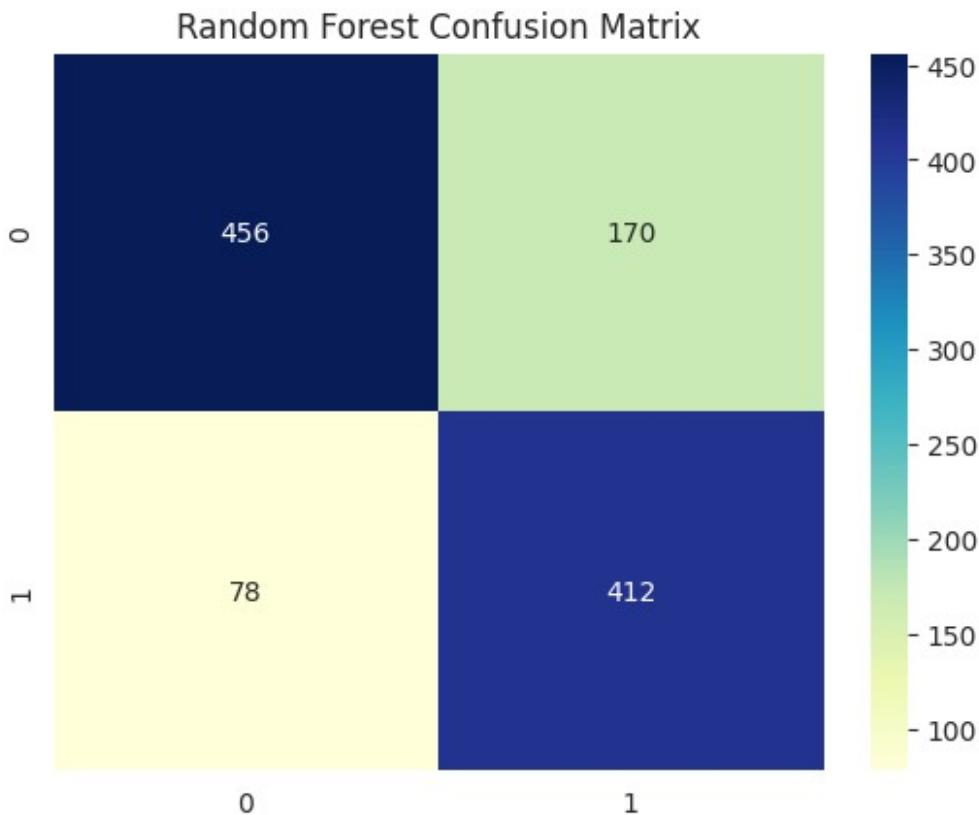
```

```

sns.heatmap(cm, annot=True, fmt='d', cmap='YlGnBu')
plt.title("Random Forest Confusion Matrix")
plt.show()

print(classification_report(y_test, rf_pred))

```



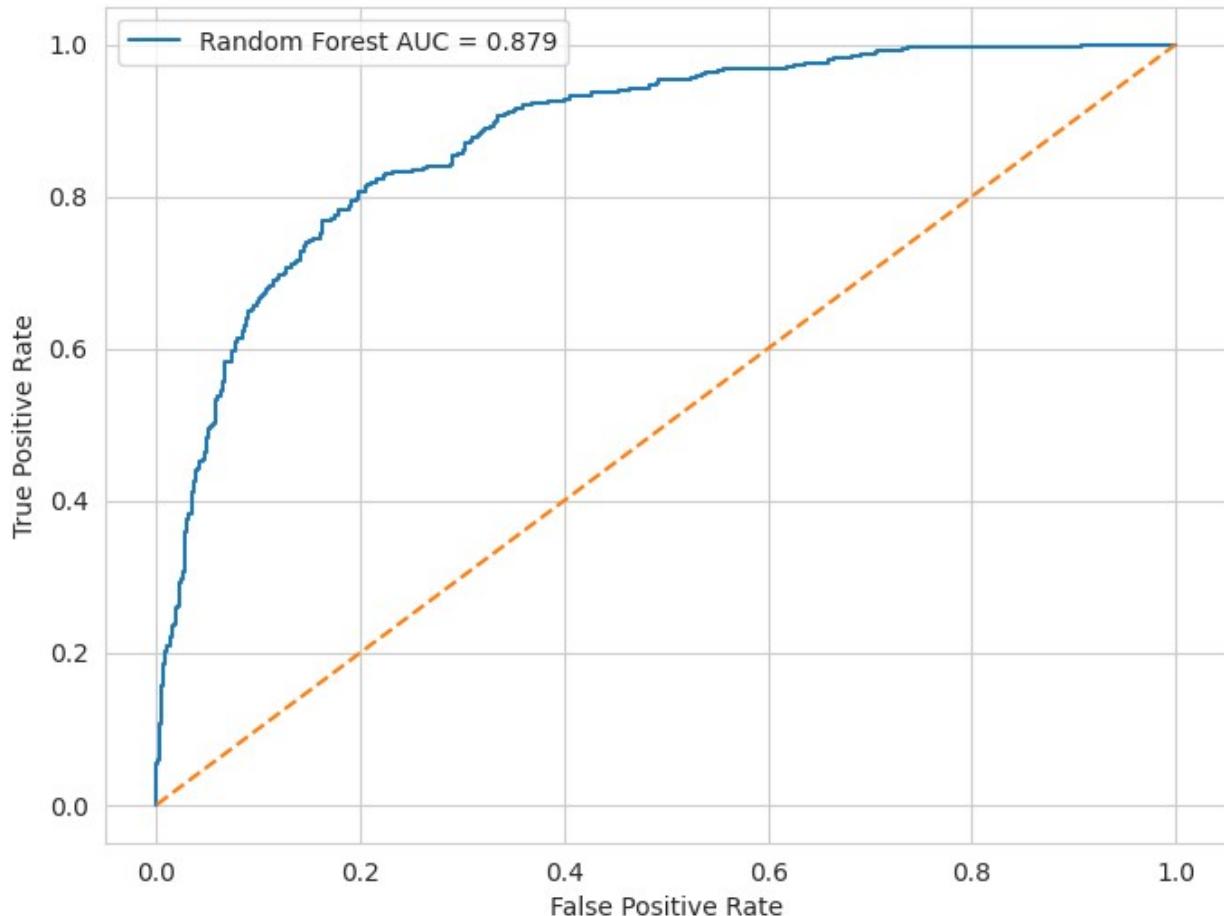
	precision	recall	f1-score	support
0	0.85	0.73	0.79	626
1	0.71	0.84	0.77	490
accuracy			0.78	1116
macro avg	0.78	0.78	0.78	1116
weighted avg	0.79	0.78	0.78	1116

```

# ROC curve and AUC
rf_probs = rf_clf.predict_proba(X_test)[:, 1]
rf_auc = roc_auc_score(y_test, rf_probs)
fpr, tpr, _ = roc_curve(y_test, rf_probs)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f"Random Forest AUC = {round(rf_auc,3)}")
plt.plot([0,1],[0,1], '--')
plt.xlabel('False Positive Rate')

```

```
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



XGBoost Classifier

```
# --- XGBoost Classifier ---
xgb_clf = xgb.XGBClassifier(
    n_estimators=200,
    learning_rate=0.05,
    max_depth=6,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42,
    eval_metric='logloss'
)
xgb_clf.fit(X_train, y_train)
xgb_pred = xgb_clf.predict(X_test)

xgb_accuracy = accuracy_score(y_test, xgb_pred)
xgb_f1 = f1_score(y_test, xgb_pred)
```

```

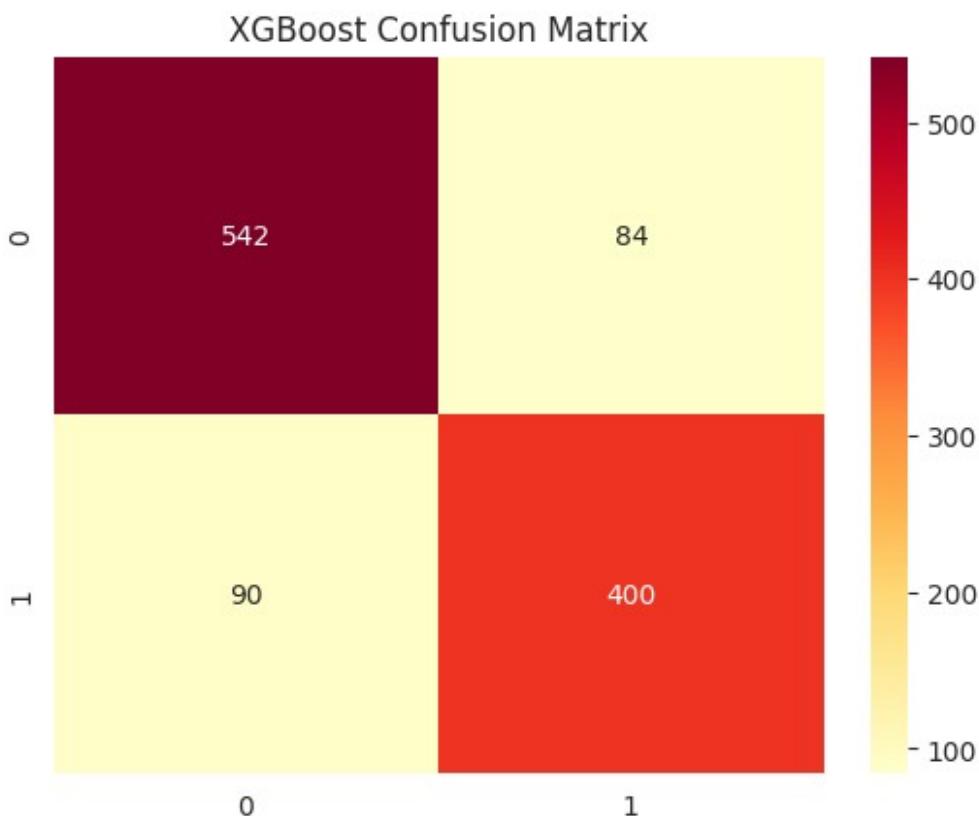
print(f"Using XGBoost we get an accuracy of
{round(xgb_accuracy*100,2)}%")
print(f"The F1 score for XGBoost is {round(xgb_f1*100,2)}%")

Using XGBoost we get an accuracy of 84.41%
The F1 score for XGBoost is 82.14%

cm = confusion_matrix(y_test, xgb_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='YlOrRd')
plt.title("XGBoost Confusion Matrix")
plt.show()

print(classification_report(y_test, xgb_pred))

```

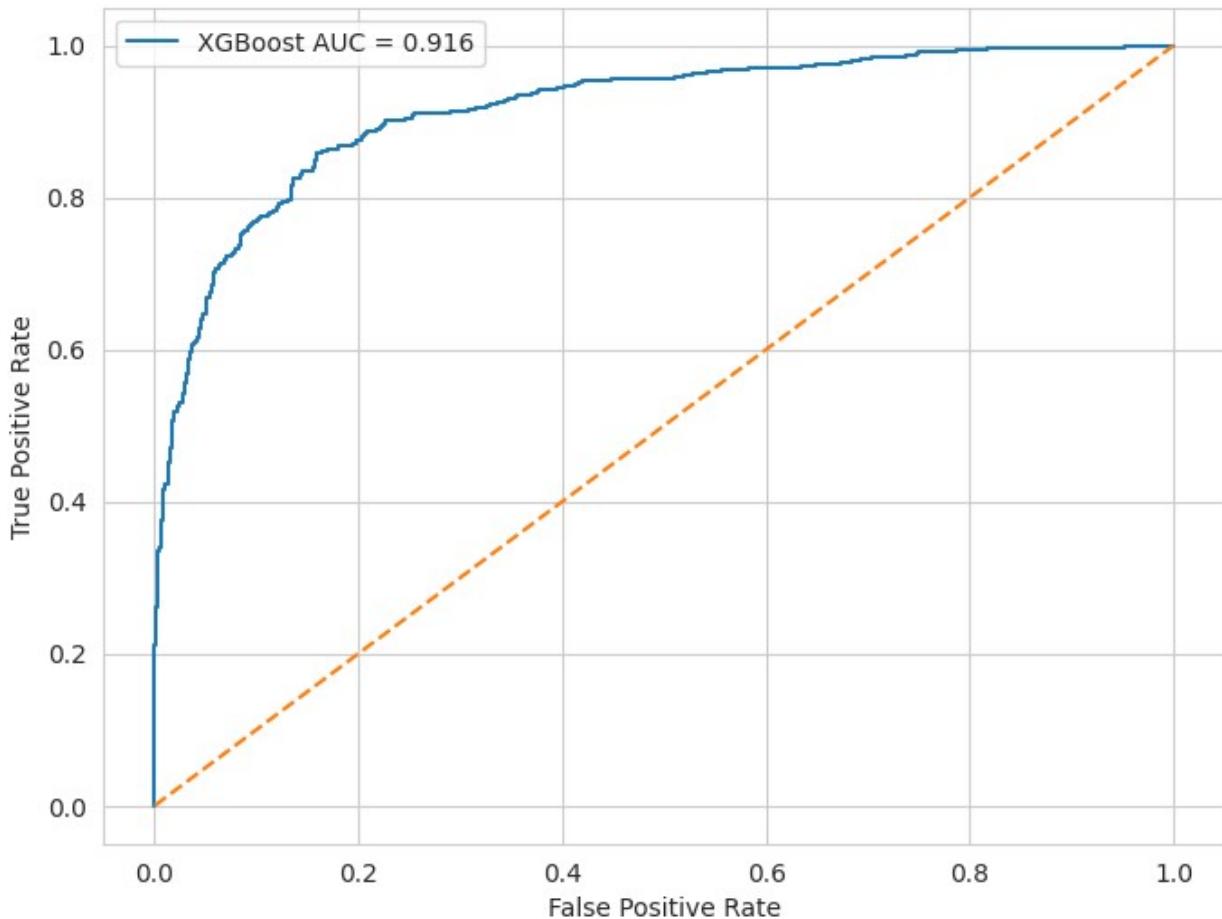


	precision	recall	f1-score	support
0	0.86	0.87	0.86	626
1	0.83	0.82	0.82	490
accuracy			0.84	1116
macro avg	0.84	0.84	0.84	1116
weighted avg	0.84	0.84	0.84	1116

```

# ROC curve and AUC
xgb_probs = xgb_clf.predict_proba(X_test)[:, 1]
xgb_auc = roc_auc_score(y_test, xgb_probs)
fpr, tpr, _ = roc_curve(y_test, xgb_probs)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f"XGBoost AUC = {round(xgb_auc,3)}")
plt.plot([0,1],[0,1], '--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```



```

comparison = pd.DataFrame({
    "Logistic regression":{'Accuracy':log_accuracy, 'AUC':log_auc, 'F1 score':logistic_f1},
    "K-nearest neighbours":{'Accuracy':knn_accuracy, 'AUC':knn_auc, 'F1 score':knn_f1},
    "Decision trees":{'Accuracy':tree_accuracy, 'AUC':tree_auc, 'F1 score':tree_f1},
    "Support vector machine":{'Accuracy':svm_accuracy, 'AUC':svm_auc,

```

```

'F1 score':svm_f1},
    "Random forest": {'Accuracy': rf_accuracy,'AUC': rf_auc, 'F1 score': rf_f1},
    "XGBoost": {'Accuracy': xgb_accuracy,'AUC': xgb_auc, 'F1 score': xgb_f1}
}).T
comparison

{"summary": {
    "name": "comparison",
    "rows": 6,
    "fields": [
        {"column": "Accuracy", "properties": {
            "dtype": "number", "std": 0.07649132373802281, "min": 0.6693548387096774, "max": 0.8772401433691757, "num_unique_values": 6, "samples": [0.6693548387096774, 0.8405017921146953, 0.8440860215053764, 0.8405017921146953, 0.8440860215053764, 0.8440860215053764]
        }, "semantic_type": "\\", "description": "\n        "},
        {"column": "AUC", "properties": {
            "dtype": "number", "std": 0.08917009810636684, "min": 0.7233650648757907, "max": 0.9483764751907153, "num_unique_values": 6, "samples": [0.7233650648757907, 0.8533937536676011, 0.9158635978353002, 0.8407871198568873, 0.8213552361396304, 0.8407871198568873]
        }, "semantic_type": "\\", "description": "\n        "},
        {"column": "F1 score", "properties": {
            "dtype": "number", "std": 0.08396792848953102, "min": 0.6448508180943214, "max": 0.8660801564027371, "num_unique_values": 6, "samples": [0.6448508180943214, 0.8213552361396304, 0.8407871198568873, 0.8213552361396304, 0.8407871198568873, 0.8213552361396304]
        }, "semantic_type": "\\", "description": "\n        "}
    ], "type": "dataframe", "variable_name": "comparison"
}

fig = plt.gcf()
fig.set_size_inches(16, 10)
titles = ['AUC', 'Accuracy', 'F1 score']

for i, label in enumerate(comparison.columns):
    plt.subplot(2, 2, i + 1)
    sns.barplot(x=comparison.index, y=comparison[label], data=comparison)
    plt.title(titles[i], fontsize=14)
    plt.ylabel('Score', fontsize=12)
    plt.xlabel('Models', fontsize=12)
    plt.xticks(rotation=25, ha='right', fontsize=10) # rotate + align
    plt.tight_layout(pad=3.0) # adds space between subplots

plt.show()

```

